



200 Kubernetes Interview Q&A

General Kubernetes Concepts

1. What is Kubernetes, and why is it important?

Answer: Kubernetes is an open-source container orchestration platform for automating deployment, scaling, and managing containerized applications. It simplifies the complexity of deploying, managing, and scaling applications, ensuring high availability, fault tolerance, and load balancing.

2. What are the main components of Kubernetes architecture?

Answer:

1. Master Node Components:

- **API Server:** Exposes Kubernetes API for interaction.
- **etcd:** Distributed key-value store for cluster data.
- **Scheduler:** Assigns workloads to nodes.
- **Controller Manager:** Manages control loops for the cluster.

2. Worker Node Components:

- **Kubelet:** Ensures containers are running in pods.
- **Kube-proxy:** Handles networking and load balancing.
- **Container Runtime:** Runs containers (e.g., Docker, CRI-O).

3. What is a Kubernetes Pod?

Answer: A pod is the smallest deployable unit in Kubernetes. It represents a single instance of a running process in a cluster, which can contain one or more tightly coupled containers that share the same network and storage.

4. How does Kubernetes handle service discovery?

Answer: Kubernetes provides service discovery through **DNS**. Each service in the cluster gets a unique DNS name (e.g., my-service.my-namespace.svc.cluster.local), and pods can communicate with services by using this DNS name.

Kubernetes Networking

5. What is a Kubernetes Service, and what are its types?

Answer: A Service is an abstraction that defines a logical set of pods and enables network access to them. Types of services include:

- **ClusterIP:** Exposes service on an internal IP within the cluster.
- **NodePort:** Exposes service on each node's IP on a static port.
- **LoadBalancer:** Exposes service externally using a cloud provider's load balancer.
- **ExternalName:** Maps a service to a DNS name outside the cluster.

6. What is kube-proxy, and what role does it play in Kubernetes?

Answer: **kube-proxy** is a network proxy that runs on each node in a Kubernetes cluster. It forwards traffic to the appropriate pod based on the service's IP address and handles load balancing for services.

7. What are Network Policies in Kubernetes?

Answer: Network Policies are used to control traffic flow at the IP address or port level within the cluster. They allow you to define which pods can communicate with which services, other pods, or external endpoints.

Pod Management and Scheduling

8. What is a DaemonSet in Kubernetes?

Answer: A **DaemonSet** ensures that a copy of a pod runs on all or some nodes in a cluster. It is commonly used for deploying node-level agents like log collectors, monitoring daemons, or network components.

9. Explain the difference between a Deployment and StatefulSet.

Answer:

- **Deployment:** Used for stateless applications and supports rolling updates, rollback, and scaling. Pods are treated as identical replicas.
- **StatefulSet:** Used for stateful applications like databases, where pods have unique identities, persistent storage, and stable network identities.

10. What is a ReplicaSet, and how does it differ from a ReplicationController?

Answer: A **ReplicaSet** ensures that a specified number of pod replicas are running at all times. It supports set-based label selectors, whereas **ReplicationController** uses equality-based selectors.

Kubernetes Storage

11. What is a PersistentVolume (PV) in Kubernetes?

Answer: A **PersistentVolume** is a piece of storage in the cluster that has been provisioned by an administrator or dynamically using a **StorageClass**. It provides a way for pods to persist data across restarts.

12. Explain the difference between PersistentVolume and PersistentVolumeClaim.

Answer:

- **PersistentVolume (PV):** Represents the physical storage resource in the cluster.
- **PersistentVolumeClaim (PVC):** A request for storage by a user that binds to an available **PersistentVolume**.

13. What is a StorageClass, and how does it relate to PersistentVolumes?

Answer: A **StorageClass** defines the types of storage (e.g., SSD, HDD, NFS) available in the cluster and allows users to dynamically provision **PersistentVolumes** with specific characteristics.

Scaling and Autoscaling

14. How does Kubernetes handle pod scaling?

Answer: Kubernetes supports horizontal scaling of pods using the **Horizontal Pod Autoscaler (HPA)**, which automatically adjusts the number of pod replicas based on CPU/memory utilization or custom metrics.

15. What is Vertical Pod Autoscaling (VPA), and how is it different from HPA?

Answer: **Vertical Pod Autoscaling (VPA)** adjusts the CPU and memory limits of a pod rather than the number of pod replicas, as in HPA. It ensures optimal resource allocation by modifying resource requests and limits over time.

Security in Kubernetes

16. What are Kubernetes Secrets, and how are they used?

Answer: Kubernetes **Secrets** are objects used to store sensitive data, such as API keys, passwords, and certificates. They ensure that sensitive information is not hardcoded in pod specifications and can be mounted as environment variables or volumes.

17. How do Role-Based Access Control (RBAC) and Roles work in Kubernetes?

Answer: **RBAC** controls who can access and perform actions within a Kubernetes cluster. **Roles** and **ClusterRoles** define sets of permissions, while **RoleBindings** and **ClusterRoleBindings** associate users or service accounts with those roles.

18. What is a ServiceAccount in Kubernetes, and when would you use it?

Answer: A **ServiceAccount** is an identity used by pods to access the Kubernetes API or other services. It is used when a pod needs to perform actions within the cluster, such as reading secrets or accessing services.

Kubernetes Troubleshooting

19. Scenario: A pod is stuck in CrashLoopBackOff. How do you troubleshoot it?

Answer:

1. Check the pod's logs using `kubectl logs <pod-name>`.
2. Use `kubectl describe pod <pod-name>` to view detailed information about the pod's status and events.
3. If the pod keeps crashing too quickly, use `kubectl exec -it <pod-name> -- /bin/sh` to get inside the running container (if it starts briefly).

20. Scenario: A pod is stuck in Pending state. What could be the issue, and how do you resolve it?

Answer:

- **Common reasons:**
 - Insufficient resources (CPU/memory).
 - NodeSelector or Affinity rules not matching any available node.
 - No PersistentVolume available for binding with a PVC.
 - **Resolution:**
 - Check `kubectl describe pod <pod-name>` for detailed events.
 - If it's a resource issue, adjust the resource limits or provision more nodes.
-

Advanced Kubernetes Concepts

21. What is a Helm Chart, and how does it simplify Kubernetes deployments?

Answer: **Helm** is a package manager for Kubernetes that enables the management of Kubernetes applications using charts. A Helm chart is a collection of files that describe a set of Kubernetes resources. It simplifies deployment, versioning, and rollback of applications.

22. What is an Ingress, and how is it used in Kubernetes?

Answer: An **Ingress** is an API object that manages external access to services in a Kubernetes cluster, typically HTTP/HTTPS traffic. It can route traffic based on hostnames, paths, or headers, providing features like SSL termination and load balancing.

23. What is a Custom Resource Definition (CRD) in Kubernetes?

Answer: A **CRD** allows you to define and create custom resources beyond the built-in Kubernetes resources. It extends the Kubernetes API to manage specific application configurations or logic as first-class objects within the cluster.

High Availability and Disaster Recovery

24. How do you ensure high availability in a Kubernetes cluster?

Answer:

1. **Multiple Master Nodes:** Deploy multiple control plane nodes across different zones.
2. **Replicated etcd:** Ensure etcd runs in a highly available, quorum-based configuration with odd numbers of replicas (3 or 5).
3. **Pod Disruption Budgets (PDBs):** Ensure that critical workloads remain available during node maintenance or scaling events.

25. What is the role of etcd in Kubernetes?

Answer: **etcd** is a distributed key-value store used by Kubernetes to store the entire cluster state, including configuration data, pod status, secrets, and service discovery information. Ensuring etcd is highly available and backed up is critical for Kubernetes disaster recovery.

Kubernetes Upgrade and Maintenance

26. How do you safely upgrade a Kubernetes cluster without downtime?

Answer:

1. Backup the etcd data using etcdctl or tools like Velero.
2. Upgrade the control plane components (API server, scheduler, controller manager) one node at a time.
3. Drain and upgrade worker nodes individually, ensuring pods are rescheduled on other nodes.
4. Monitor the cluster's health using `kubectl get nodes` and `kubectl get pods`.

27. What are taints and tolerations in Kubernetes?

Answer: **Taints** allow nodes to repel certain pods unless those pods have matching **tolerations**. This ensures that only specific pods are scheduled on particular nodes (e.g., dedicated nodes for system workloads or specialized hardware).

Kubernetes Monitoring and Logging

28. How do you monitor Kubernetes clusters and applications?

Answer: Common tools for Kubernetes monitoring include:

- **Prometheus:** Collects and stores metrics.
- **Grafana:** Visualizes metrics from Prometheus in customizable dashboards.
- **Kubernetes Metrics Server:** Provides resource usage metrics for HPA.

29. What is the EFK stack, and how is it used for logging?

Answer: The **EFK stack** consists of **Elasticsearch**, **Fluentd**, and **Kibana**. Fluentd collects logs from Kubernetes pods and forwards them to Elasticsearch, where they are stored and indexed. Kibana is used to visualize and analyze the logs.

Advanced Scenario-Based Questions

30. Scenario: You need to implement canary deployment with Kubernetes and Istio. How do you approach it?

Answer:

1. Deploy both versions (v1 and v2) of the service.
2. Define an Istio **VirtualService** to route 90% of the traffic to v1 and 10% to v2 initially.
3. Gradually shift more traffic to v2 while monitoring for errors.
4. If successful, move 100% of the traffic to v2 and retire v1.

31. Scenario: You are asked to scale an application that spikes in traffic during certain hours. What steps do you take to ensure autoscaling?

Answer:

1. **Enable Horizontal Pod Autoscaling (HPA):**
 - Define resource requests and limits in the pod's configuration.
 - Deploy **HPA** based on CPU, memory, or custom metrics.

```
kubectl autoscale deployment <deployment-name> --cpu-percent=60 --min=2 --max=10
```

2. **Monitor Metrics:** Use **Prometheus** or other monitoring tools to ensure HPA is working correctly.
3. **Vertical Pod Autoscaler (VPA):** Use VPA if the bottleneck is with CPU/memory limits on individual pods.

32. Scenario: Your pods cannot connect to an external service. How do you troubleshoot this issue?

Answer:

1. Check if the pods have external internet access using `kubectl exec <pod-name> -- curl <external-service>`.
2. Verify network policies to ensure they are not blocking outbound traffic.
3. Check service configurations (ClusterIP, NodePort, or LoadBalancer) to verify external service exposure.
4. Ensure **kube-proxy** is running and properly routing traffic between pods and external services.

33. Scenario: You want to restrict communication between pods of different namespaces. How do you enforce it?

Answer: Use **Network Policies** to control traffic between namespaces.

1. Create a network policy that allows ingress traffic only from pods within the same namespace:

apiVersion: networking.k8s.io/v1

kind: NetworkPolicy

metadata:

name: restrict-cross-namespace

namespace: app-namespace

spec:

podSelector: {}

policyTypes:

- Ingress

ingress:

- from:

- podSelector: {}

34. Scenario: Your Kubernetes cluster is running out of disk space due to large container logs. How do you fix this issue?

Answer:

1. **Log Rotation:** Implement log rotation in the container runtime (Docker, CRI-O). For Docker, update `/etc/docker/daemon.json`:

```
{  
  "log-driver": "json-file",  
  "log-opts": {  
    "max-size": "100m",  
    "max-file": "3"  
  }  
}
```

2. **Centralized Logging:** Offload logs to a centralized logging solution like EFK (Elasticsearch, Fluentd, Kibana) or Loki/Promtail.
3. **Monitor Disk Usage:** Use Prometheus and Grafana to track disk usage over time.

35. Scenario: You need to roll back a failed deployment. What are the steps to revert to a previous version?

Answer:

1. Check deployment history:

```
kubectl rollout history deployment <deployment-name>
```

2. Roll back to the previous stable version:

```
kubectl rollout undo deployment <deployment-name>
```

3. Verify the deployment status:

```
kubectl get pods -w
```

36. Scenario: You are tasked with deploying a stateful application (like a database) in Kubernetes. What approach do you take?

Answer:

1. Use a **StatefulSet** to deploy the application, as StatefulSets ensure stable network identities and persistent storage.
2. Provision a **PersistentVolumeClaim** (PVC) for each pod to ensure data persistence.
3. Set up a **Headless Service** to maintain the network identity of each replica.

Example for MySQL deployment:

apiVersion: apps/v1

kind: StatefulSet

metadata:

name: mysql

spec:

serviceName: "mysql"

replicas: 3

template:

spec:

containers:

- name: mysql

image: mysql:5.7

volumeMounts:

- name: mysql-persistent-storage

mountPath: /var/lib/mysql

volumeClaimTemplates:

```
- metadata:  
  name: mysql-persistent-storage  
spec:  
  accessModes: ["ReadWriteOnce"]  
resources:  
  requests:  
    storage: 1Gi
```

37. Scenario: A Kubernetes node is unresponsive, and you need to migrate the pods to a different node. What steps do you take?

Answer:

1. Drain the Node:

```
kubectldrain <node-name> --ignore-daemonsets
```

2. Cordon the Node to prevent new pods from being scheduled:

```
kubectlcordon <node-name>
```

3. Pods running on the node will be rescheduled on other available nodes.

Monitor the migration process:

```
kubectlget pods -o wide
```

Kubernetes Security

38. What are Pod Security Policies (PSP) in Kubernetes?

Answer: Pod Security Policies are cluster-level resources that control security-sensitive aspects of pod specification, such as:

- Whether a pod can run as privileged.
- Restricting container capabilities.

- Enforcing the use of specific security contexts, such as running as a non-root user.

39. How do you secure sensitive data in Kubernetes using Secrets?

Answer:

1. Create a Secret:

apiVersion: v1

kind: Secret

metadata:

name: db-credentials

type: Opaque

data:

username: YWRtaW4= # base64 encoded 'admin'

password: cGFzc3dvcmQ= # base64 encoded 'password'

2. Mount the Secret as an environment variable or volume:

env:

- name: DB_USERNAME

valueFrom:

secretKeyRef:

name: db-credentials

key: username

40. What is Kubernetes Role-Based Access Control (RBAC)?

Answer: RBAC in Kubernetes is a way to control who can perform specific actions within the cluster. It defines roles (sets of permissions) and associates those roles with users, groups, or service accounts using **RoleBindings** or **ClusterRoleBindings**.

Kubernetes Ingress and Traffic Management

41. What is Kubernetes Ingress, and how is it different from a service?

Answer: An **Ingress** is an API object that manages external HTTP/HTTPS access to services, typically used for routing traffic to different services based on hostnames or paths. Unlike a service, which exposes individual pods or deployments, an Ingress acts as a gateway that controls and manages multiple backend services.

42. Scenario: You need to set up SSL termination for your services in Kubernetes. How do you configure it using Ingress?

Answer:

1. Generate or obtain an SSL certificate (can be done using **cert-manager** and Let's Encrypt).
2. Create a Kubernetes **Secret** to store the SSL certificate:

```
kubectl create secret tls tls-secret --cert=path/to/tls.crt --key=path/to/tls.key
```

3. Define an Ingress resource with TLS enabled:

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: myapp-ingress
```

```
spec:
```

```
  tls:
```

```
  - hosts:
```

```
    - myapp.example.com
```

```
    secretName: tls-secret
```

```
  rules:
```

- host: myapp.example.com

http:

paths:

- path: /

pathType: Prefix

backend:

service:

name: myapp-service

port:

number: 80

Kubernetes Deployment and Scaling

43. What is the purpose of Kubernetes Deployments?

Answer: Kubernetes **Deployments** provide declarative updates to applications. They manage the lifecycle of applications by:

- Creating new replica sets and scaling them.
- Performing rolling updates or rollbacks.
- Monitoring the health of applications and rescheduling failed pods.

44. What is the difference between a StatefulSet and a DaemonSet?

Answer:

- **StatefulSet:** Used for stateful applications where each pod has a unique identity (e.g., databases, distributed systems).
- **DaemonSet:** Ensures that a pod runs on every node (or a subset of nodes), typically used for logging, monitoring, or networking components.

Kubernetes Monitoring and Logging

45. How do you monitor a Kubernetes cluster?

Answer: Common tools for monitoring include:

- **Prometheus:** Collects metrics from Kubernetes components and applications.
- **Grafana:** Visualizes metrics from Prometheus or other sources in real-time dashboards.
- **Kubernetes Metrics Server:** Provides resource usage metrics (CPU, memory) for pods and nodes, used by the Horizontal Pod Autoscaler.

46. Scenario: You need to centralize logs from all pods in the cluster. How do you achieve this?

Answer:

1. Deploy a **logging agent** like Fluentd or Fluent Bit on each node as a DaemonSet.
2. Set up a **centralized logging solution** like Elasticsearch or Loki to collect logs from Fluentd.
3. Use **Kibana** or **Grafana** for log visualization and analysis.

Kubernetes High Availability and Disaster Recovery

47. How do you ensure High Availability (HA) for the Kubernetes control plane?

Answer:

1. Run multiple control plane nodes across different availability zones.
2. Ensure **etcd** is highly available with 3 or 5 members distributed across zones.
3. Use a **load balancer** to distribute traffic to multiple API server instances.

4. Implement **Pod Disruption Budgets (PDB)** for critical components to ensure that minimum replicas remain available during disruptions.

48. Scenario: The etcd database has crashed. How do you recover it?

Answer:

1. Restore the **etcd** backup (using **etcdctl** or a backup tool like Velero):

```
etcdctl snapshot restore <snapshot.db> --data-dir=<restore-dir>
```

2. Restart the etcd cluster with the restored data.
3. Ensure that the restored etcd data is consistent with the current state of the Kubernetes cluster by verifying the cluster health using `kubectl get nodes/pods`.

Kubernetes Upgrades and Maintenance

49. How do you perform a rolling update in Kubernetes?

Answer: A rolling update can be done by modifying the deployment:

1. Update the container image in the deployment YAML or using `kubectl`:

```
kubectl set image deployment/myapp mycontainer=myimage:v2
```

2. Monitor the rolling update status:

```
kubectl rollout status deployment/myapp
```

50. What is a Pod Disruption Budget (PDB), and why is it important?

Answer: A **Pod Disruption Budget** specifies the minimum number or percentage of pods that must remain available during voluntary disruptions (e.g., node upgrades, scaling events). It helps prevent critical workloads from being disrupted during maintenance operations.

Advanced Kubernetes Features and Scenarios

51. Scenario: A critical service in your application needs to be updated without any downtime. How do you ensure zero downtime during deployment?

Answer:

1. **Rolling Update:** Kubernetes supports rolling updates natively. Update the deployment by changing the container image or configuration.

- Example command:

```
kubectl set image deployment/myapp mycontainer=myimage:v2
```

2. **Readiness Probes:** Configure readiness probes to ensure that traffic is only sent to pods that are fully initialized.

- Add readiness probe to the pod:

readinessProbe:

httpGet:

path: /health

port: 8080

initialDelaySeconds: 10

periodSeconds: 5

3. **Monitor Rollout:** Use `kubectl rollout status` to monitor the progress of the deployment.

52. Scenario: You need to expose multiple services behind a single IP address using path-based routing. How do you achieve this?

Answer:

1. Use **Ingress** to configure path-based routing.
2. Example Ingress resource:

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: multi-service-ingress
```

spec:

rules:

- host: "example.com"

http:

paths:

- path: /app1

pathType: Prefix

backend:

service:

name: app1-service

port:

number: 80

- path: /app2

pathType: Prefix

backend:

service:

name: app2-service

port:

number: 80

3. This will route traffic to /app1 and /app2 based on the path.

53. Scenario: Your pods are being evicted frequently. How do you identify and resolve the issue?

Answer:

1. **Check Node Status:** Inspect the node's resources using:

kubectl describe node <node-name>

2. Investigate Pod Evictions:

- Look at pod events:

kubectl describe pod <pod-name>

- Reasons could be related to **resource pressure** (CPU, memory), disk pressure, or node conditions.

3. Use Resource Quotas and Limits: Ensure that your pods have defined resource requests and limits to prevent overconsumption.

resources:

requests:

memory: "100Mi"

cpu: "250m"

limits:

memory: "200Mi"

cpu: "500m"

4. Vertical Pod Autoscaler (VPA): Consider using VPA to adjust pod resource limits dynamically.

54. Scenario: You want to allow an application to run only on nodes with specific hardware (e.g., GPU). How do you configure node selection?

Answer:

1. Label the nodes:

kubectl label nodes <node-name> hardware=gpu

2. Use NodeSelector in the pod specification to schedule the pod on the labeled nodes:

spec:

nodeSelector:

hardware: gpu

55. Scenario: A Kubernetes node is overloaded with CPU/memory utilization. How do you troubleshoot and resolve it?

Answer:

1. **Check node resource utilization** using `kubectl top nodes`.
2. **Identify resource-heavy pods** using:

`kubectl top pod --all-namespaces`

3. **Vertical Pod Autoscaler (VPA)**: Implement VPA to manage resource limits dynamically.
4. **Reschedule pods** by draining the overloaded node:

`kubectl drain <node-name> --ignore-daemonsets --force`

56. Scenario: Your application deployment requires a specific sequence of starting and stopping pods. How do you ensure correct pod startup order?

Answer:

1. Use **Init Containers**: Ensure some actions are performed before the main containers start.
 - Example of an init container:

initContainers:

- name: init-myservice

image: busybox

command: ['sh', '-c', 'echo Initializing Service']

2. For startup ordering between different pods, use **StatefulSet** for stateful services where pod ordering is critical. StatefulSets ensure ordered, stable deployment.

Kubernetes Security

57. Scenario: You want to limit the permissions for a specific pod to access Kubernetes API. How do you enforce this?

Answer:

1. **Service Account:** Create a service account with minimal access.

```
kubectl create serviceaccount my-serviceaccount
```

2. **Role and RoleBinding:**

- Create a role with restricted permissions:

```
kind: Role
```

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
metadata:
```

```
  namespace: default
```

```
  name: restricted-role
```

```
rules:
```

```
- apiGroups: [""]
```

```
  resources: ["pods"]
```

```
  verbs: ["get", "list"]
```

- Bind the role to the service account:

```
kubectl create rolebinding my-binding --role=restricted-role --  
serviceaccount=default:my-serviceaccount
```

58. Scenario: You need to ensure that a container runs without root privileges. How do you enforce this?

Answer:

1. Set the securityContext in the pod to ensure non-root execution:

securityContext:

runAsUser: 1000

runAsGroup: 3000

fsGroup: 2000

2. **PodSecurityPolicies (PSP)** can be used to enforce security policies across the cluster. Set policies like MustRunAsNonRoot.
-

Kubernetes Scaling

59. How does Kubernetes Horizontal Pod Autoscaler (HPA) work?

Answer:

1. HPA automatically scales the number of pod replicas based on resource utilization (e.g., CPU, memory).
2. It adjusts the number of pods using metrics provided by the Metrics API (default metrics include CPU and memory).
 - Example HPA configuration:

```
kubectl autoscale deployment myapp --cpu-percent=50 --min=2 --max=10
```

60. Scenario: You want to scale your pods based on custom application metrics (e.g., requests per second). How do you achieve this?

Answer:

1. **Expose Custom Metrics:** Use Prometheus or a similar tool to expose custom metrics.

2. **Install Custom Metrics Adapter:** Install Prometheus Adapter to expose custom metrics in Kubernetes.

3. **Configure HPA with custom metrics:**

apiVersion: autoscaling/v2beta2

kind: HorizontalPodAutoscaler

metadata:

name: custom-hpa

spec:

scaleTargetRef:

apiVersion: apps/v1

kind: Deployment

name: myapp

minReplicas: 1

maxReplicas: 10

metrics:

- type: Pods

pods:

metric:

name: http_requests_per_second

target:

type: AverageValue

averageValue: "100"

Kubernetes Upgrades and Maintenance

61. Scenario: You need to upgrade Kubernetes to a new version without disrupting your services. What steps do you follow?

Answer:

1. Backup etcd:

- Ensure you have a backup of etcd before performing any upgrades.

`etcdctl snapshot save snapshot.db`

2. Upgrade Control Plane Components:

- Upgrade control plane components (API server, scheduler, etc.) one node at a time.

`kubeadm upgrade plan`

`kubeadm upgrade apply v1.x.x`

3. Drain and Upgrade Worker Nodes:

- Drain one node at a time, upgrade kubelet, and uncordon:

`kubectldrain <node-name> --ignore-daemonsets`

`apt-get update && apt-get install -y kubelet=1.x.x-00`

`kubectldrain <node-name>`

62. Scenario: During a node maintenance event, you want to minimize disruption to critical pods. What features in Kubernetes help achieve this?

Answer:

- 1. Pod Disruption Budgets (PDB):** Define PDBs to ensure a minimum number of pods remain available during voluntary disruptions.

- Example PDB configuration:

```
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: myapp-pdb
spec:
  minAvailable: 80%
  selector:
    matchLabels:
      app: myapp
```

Kubernetes Monitoring and Troubleshooting

63. Scenario: A pod is stuck in CrashLoopBackOff. How do you troubleshoot it?

Answer:

1. **Check pod logs:**

kubectl logs <pod-name>

2. **Describe the pod** to check events:

kubectl describe pod <pod-name>

3. **Investigate Liveness/Readiness Probes:** Misconfigured probes can cause pod restarts.

4. Check for **resource constraints** (memory or CPU) that might be causing the container to terminate.

64. Scenario: A service is not accessible from outside the cluster, even though a NodePort or LoadBalancer service is configured. How do you troubleshoot this?

Answer:

1. **Check Service Configuration:**

```
kubectl get svc <service-name> -o yaml
```

2. **Ensure NodePort or LoadBalancer IP is accessible.** For LoadBalancer, ensure the cloud provider has provisioned the external load balancer.
3. **Check Network Policies:** Ensure that ingress traffic is allowed if network policies are enforced.
4. Verify that **kube-proxy** is running and forwarding traffic correctly.

Kubernetes Advanced Scenarios

65. Scenario: You need to deploy an application across multiple clusters with consistent configuration. How do you achieve this?

Answer:

1. **KubeFed (Kubernetes Federation):** Use **KubeFed** to manage multiple clusters from a single control plane.
 - Install KubeFed:

```
kubectl apply -f https://github.com/kubernetes-sigs/kubefed/releases/download/v0.1.0/kubefed.yaml
```

2. **Use Helm:** Define the application using Helm charts and deploy them consistently across all clusters.
3. **CI/CD pipelines:** Integrate with Jenkins or GitLab for continuous deployment across clusters.

66. Scenario: You are tasked with deploying a service mesh to manage communication between microservices. How do you set up Istio in Kubernetes?

Answer:

1. **Install Istio:**
 - Use Istioctl to install the service mesh.

```
istioctl install --set profile=default
```

2. **Enable automatic sidecar injection** for your namespace:

kubectl label namespace default istio-injection=enabled

3. **Traffic Management:** Use Istio's VirtualService and DestinationRule to control traffic between microservices:

```
apiVersion: networking.istio.io/v1alpha3
```

```
kind: VirtualService
```

```
metadata:
```

```
  name: myapp-route
```

```
spec:
```

```
  hosts:
```

```
  - myapp.example.com
```

```
  http:
```

```
  - route:
```

```
    - destination:
```

```
      host: myapp
```

```
      subset: v1
```

67. Scenario: Pods in your cluster are experiencing DNS resolution issues. How do you troubleshoot and resolve the issue?

Answer:

1. **Check DNS ConfigMap:** Verify that the CoreDNS ConfigMap is correctly configured:

```
kubectl -n kube-system get configmap coredns -o yaml
```

2. **DNS Resolution Test:** Check DNS resolution inside a pod using nslookup or dig:

```
kubectl exec -it <pod-name> -- nslookup kubernetes.default
```

3. **Verify CoreDNS Deployment:** Ensure CoreDNS pods are running correctly:

```
kubectl -n kube-system get pods -l k8s-app=kube-dns
```

4. **Network Policies:** Ensure that network policies are not blocking DNS traffic (port 53).

68. Scenario: A pod is stuck in ContainerCreating state. How do you troubleshoot and resolve it?

Answer:

1. **Check Events:** Describe the pod to see why it is stuck:

```
kubectl describe pod <pod-name>
```

2. **Image Pull Issues:** If the issue is related to pulling the container image, check:
 - Correct image name, tag, and registry credentials.
 - Network connectivity to the image registry.
3. **Check node status:** Ensure the node has sufficient resources (CPU, memory, disk space) to create the container.
4. **Persistent Volume Issues:** If the pod is waiting for a volume to be attached, verify that the PersistentVolume is correctly bound.

69. Scenario: Your service is exposed using a LoadBalancer, but it's not receiving external traffic. How do you troubleshoot this?

Answer:

1. **Check Service Type:** Verify that the service is of type LoadBalancer:

```
kubectl get svc <service-name> -o yaml
```

2. **Cloud Provider Configuration:** Ensure that the cloud provider has created an external load balancer (for AWS, GCP, Azure).
 - Check the external IP of the LoadBalancer using `kubectl get svc`.

3. **Firewall Rules:** Ensure the external firewall (security groups, network ACLs) allows traffic on the LoadBalancer ports.
 4. **Check kube-proxy:** Ensure kube-proxy is functioning correctly and forwarding traffic to the correct pods.
-

Kubernetes Ingress and Traffic Management

70. Scenario: Your Ingress is not routing traffic correctly to the backend services. What steps do you take to troubleshoot this?

Answer:

1. **Check Ingress Configuration:** Ensure that the Ingress rules are correctly configured:

```
kubectl get ingress <ingress-name> -o yaml
```

2. **Verify Service Endpoints:** Ensure that the service behind the Ingress is healthy and has active endpoints:

```
kubectl get endpoints <service-name>
```

3. **Ingress Controller:** Verify that the Ingress controller (e.g., NGINX, Traefik) is running and healthy:

```
kubectl get pods -n ingress-nginx
```

4. **DNS Resolution:** Ensure that the DNS for the Ingress host (e.g., example.com) resolves to the correct external IP.

71. Scenario: You need to set up a canary deployment to shift 10% of traffic to a new version of the service. How do you configure it?

Answer:

1. Deploy two versions of the service (v1 and v2).
2. Use **Istio** or **NGINX Ingress** to split traffic between the two versions:
 - Example using Istio:

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

name: myapp

spec:

hosts:

- myapp.example.com

http:

- route:

- destination:

- host: myapp

- subset: v1

- weight: 90

- destination:

- host: myapp

- subset: v2

- weight: 10

Kubernetes Storage and Persistent Volumes

72. Scenario: A pod is unable to mount a PersistentVolume. What steps do you take to resolve this?

Answer:

1. **Check PVC Binding:** Ensure that the PersistentVolumeClaim (PVC) is correctly bound to a PersistentVolume (PV):

`kubectl get pvc <pvc-name>`

2. **Describe Events:** Use `kubectl describe pod <pod-name>` to view events related to volume mounting.
3. **Verify StorageClass:** If dynamic provisioning is used, ensure that the StorageClass is correctly defined and provisioned.
4. **Check Node Status:** Ensure the node has access to the storage backend (NFS, EBS, etc.).

73. Scenario: You need to resize a PersistentVolume that is dynamically provisioned. How do you handle this?

Answer:

1. Ensure that the underlying storage class supports volume resizing (`allowVolumeExpansion: true`).
2. Edit the PVC to request a larger size:

`apiVersion: v1`

`kind: PersistentVolumeClaim`

`spec:`

`resources:`

`requests:`

`storage: 10Gi`

3. Verify that the volume is resized and the new capacity is reflected using:

`kubectl describe pvc <pvc-name>`

74. Scenario: Your application requires a shared persistent volume between multiple pods. How do you configure this in Kubernetes?

Answer:

1. Use a PersistentVolume that supports shared access modes, such as **ReadWriteMany (RWX)**.

2. Example of NFS-based shared volume:

apiVersion: v1

kind: PersistentVolume

metadata:

name: shared-pv

spec:

capacity:

storage: 10Gi

accessModes:

- ReadWriteMany

nfs:

path: /exported/path

server: nfs-server.example.com

Kubernetes High Availability and Scaling

75. Scenario: You want to deploy a highly available database cluster in Kubernetes. What approach do you take?

Answer:

1. **Use StatefulSets:** Deploy the database using a **StatefulSet**, which ensures stable network identities and persistent storage for each replica.

- Example StatefulSet for MySQL:

apiVersion: apps/v1

kind: StatefulSet

metadata:

name: mysql

spec:

serviceName: "mysql"

replicas: 3

template:

spec:

containers:

- name: mysql

image: mysql:5.7

volumeMounts:

- name: mysql-persistent-storage

mountPath: /var/lib/mysql

volumeClaimTemplates:

- metadata:

name: mysql-persistent-storage

spec:

accessModes: ["ReadWriteOnce"]

resources:

requests:

storage: 10Gi

2. **Headless Service:** Ensure you use a **headless service** to allow for stable DNS entries for each replica.
3. **Configure replication:** Configure MySQL or any other database for replication between nodes.

76. Scenario: Your cluster is running out of resources (CPU, memory). How do you scale the nodes in your cluster automatically?

Answer:

1. **Cluster Autoscaler:** Install and configure **Cluster Autoscaler** to automatically scale nodes based on resource usage.

- For GKE:

```
gcloud container clusters update my-cluster --enable-autoscaling --min-nodes=3 --max-nodes=10
```

- For self-managed clusters, deploy the Cluster Autoscaler using YAML configurations.

2. Ensure your pods have defined CPU and memory resource requests, which will trigger the autoscaler when the cluster is under pressure.

Kubernetes CI/CD Integration

77. Scenario: You need to implement a CI/CD pipeline that deploys a new version of an application to Kubernetes whenever a new code change is pushed to Git. How do you achieve this?

Answer:

1. **Jenkins Pipeline:** Set up a Jenkins pipeline with the Kubernetes plugin.

- Define a pipeline in .Jenkinsfile:

```
pipeline {
  agent any
  stages {
    stage('Build') {
      steps {
        sh 'docker build -t myapp:latest .'
        sh 'docker push myapp:latest'
      }
    }
  }
}
```

```
}  
stage('Deploy') {  
  steps {  
    kubernetesDeploy(  
      configs: 'deployment.yaml',  
      kubeconfigId: 'kubeconfig'  
    )  
  }  
}  
}
```

2. **GitLab CI:** Alternatively, use GitLab CI's Kubernetes integration to trigger deployments on code push.

Kubernetes Logging and Monitoring

78. Scenario: You want to collect logs from all pods and centralize them in a single location. How do you achieve this?

Answer:

1. Deploy a **Fluentd** or **Fluent Bit** DaemonSet on each node to collect logs from all containers.
2. Forward logs to a centralized logging system, such as **Elasticsearch** or **Loki**.
3. Use **Kibana** or **Grafana** to visualize and analyze the logs.

79. Scenario: Your cluster is experiencing high CPU and memory utilization, and you need to identify the bottleneck. How do you troubleshoot?

Answer:

1. Use `kubectl top nodes` and `kubectl top pods` to view resource usage.
 2. **Prometheus:** Set up Prometheus to collect detailed metrics from the cluster.
 3. **Grafana:** Use Grafana to visualize metrics such as CPU, memory, and network traffic.
 4. Identify specific pods or services consuming excessive resources and adjust resource limits or scale accordingly.
-

Kubernetes RBAC and Security

80. Scenario: You need to restrict access to a particular namespace so only a specific team can make changes. How do you configure this?

Answer:

1. **Create Role and RoleBinding:**

- Define a Role with permissions for the team:

kind: Role

apiVersion: rbac.authorization.k8s.io/v1

metadata:

namespace: dev

name: dev-role

rules:

- apiGroups: [""]

resources: ["pods", "services"]

verbs: ["get", "list", "create", "delete"]

- Create a RoleBinding to associate the Role with the team's users or service accounts:

```
kubectl create rolebinding dev-binding --role=dev-role --user=team-member
```

81. Scenario: You want to enforce security best practices by ensuring that no pod can run with root privileges. How do you implement this?

Answer:

1. **PodSecurityPolicies:** Create a PodSecurityPolicy that prevents pods from running as root:

```
apiVersion: policy/v1beta1
```

```
kind: PodSecurityPolicy
```

```
metadata:
```

```
  name: restricted
```

```
spec:
```

```
  privileged: false
```

```
  runAsUser:
```

```
    rule: MustRunAsNonRoot
```

```
  seLinux:
```

```
    rule: RunAsAny
```

```
  fsGroup:
```

```
    rule: MustRunAs
```

2. **Apply PodSecurityPolicy** with RoleBindings to enforce it across the cluster.

82. Scenario: You need to implement Role-Based Access Control (RBAC) for a specific Kubernetes namespace where developers can only read pod logs but cannot modify any resources. How do you achieve this?

Answer:

1. **Create a Role** that allows reading pods and accessing logs:

DevOps Shack

kind: Role

apiVersion: rbac.authorization.k8s.io/v1

metadata:

namespace: dev

name: pod-log-reader

rules:

- apiGroups: [""]

resources: ["pods/log"]

verbs: ["get", "list"]

2. Bind the Role to Users or Groups:

```
kubectl create rolebinding log-reader-binding --role=pod-log-reader --user=developer
```

This restricts developers to only reading logs in the dev namespace.

83. Scenario: How do you restrict certain Kubernetes workloads from running as privileged containers?

Answer:

1. PodSecurityPolicy (PSP): Enforce the policy to deny privileged containers:

apiVersion: policy/v1beta1

kind: PodSecurityPolicy

metadata:

name: restricted

spec:

privileged: false

runAsUser:

rule: MustRunAsNonRoot

DevOps Shack

seLinux:

rule: RunAsAny

fsGroup:

rule: MustRunAs

volumes:

- configMap

- secret

2. Bind this policy to the cluster or namespace using **RoleBindings** or **ClusterRoleBindings** to apply it to workloads.

Kubernetes Network Policies

84. Scenario: You need to isolate two namespaces so that pods in one namespace cannot communicate with pods in another namespace. How do you configure this?

Answer:

1. **NetworkPolicy:** Define network policies for both namespaces to block ingress traffic from the other namespace.
 - o Example for namespace dev:

apiVersion: networking.k8s.io/v1

kind: NetworkPolicy

metadata:

name: deny-cross-namespace-traffic

namespace: dev

spec:


```
podSelector: {}  
policyTypes:  
- Ingress  
ingress:  
- from:  
  - namespaceSelector:  
    matchLabels:  
      name: dev
```

This ensures that only traffic within the same namespace (dev) is allowed.

85. Scenario: Your application has strict regulatory compliance requirements, and you need to control which IP ranges can access the application. How do you implement this in Kubernetes?

Answer:

1. **NetworkPolicy:** Define a network policy that restricts ingress traffic to specific IP ranges:

```
apiVersion: networking.k8s.io/v1  
kind: NetworkPolicy  
metadata:  
  name: restrict-ingress  
  namespace: prod  
spec:  
  podSelector: {}  
  policyTypes:  
  - Ingress  
  ingress:
```

- from:

- ipBlock:

cidr: 192.168.1.0/24

except:

- 192.168.1.100/32

2. This allows traffic only from the specified IP block, while excluding any addresses you don't want to allow.

Kubernetes Multi-Cluster Management

86. Scenario: Your organization manages multiple Kubernetes clusters across different cloud providers. How do you ensure consistent deployment and management across these clusters?

Answer:

1. **KubeFed (Kubernetes Federation):** Use KubeFed to manage multiple clusters from a central control plane.

- Install KubeFed to manage multi-cluster deployments:

```
kubectl apply -f https://github.com/kubernetes-sigs/kubefed/releases/download/v0.1.0/kubefed.yaml
```

2. **Helm:** Package your application as a Helm chart for consistent deployment across clusters.
3. Use **GitOps** tools like **ArgoCD** or **Flux** to manage deployments across clusters with a single source of truth (Git repository).

Kubernetes Storage and Persistent Volumes

87. Scenario: You want to dynamically provision storage for your application, but each environment (dev, prod) uses different storage backends. How do you handle this?

Answer:

1. Use **StorageClasses** to define different storage backends for each environment:

- For production (e.g., AWS EBS):

apiVersion: storage.k8s.io/v1

kind: StorageClass

metadata:

name: prod-storage

provisioner: kubernetes.io/aws-ebs

parameters:

type: gp2

- For development (e.g., NFS):

apiVersion: storage.k8s.io/v1

kind: StorageClass

metadata:

name: dev-storage

provisioner: nfs

2. When creating PersistentVolumeClaims, specify the appropriate StorageClass:

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

name: dev-pvc

spec:

storageClassName: dev-storage

accessModes:

- ReadWriteOnce

resources:

requests:

storage: 5Gi

88. Scenario: You need to ensure that all data stored in PersistentVolumes is encrypted. How do you configure this in Kubernetes?

Answer:

1. Use a cloud provider's encrypted storage (e.g., AWS EBS with encryption enabled).
 - For AWS, ensure the EBS volume is encrypted at rest by specifying the encrypted parameter in the StorageClass:

apiVersion: storage.k8s.io/v1

kind: StorageClass

metadata:

name: encrypted-ebs

provisioner: kubernetes.io/aws-ebs

parameters:

encrypted: "true"

type: gp2

89. Scenario: Your application requires high availability and persistent storage that spans multiple availability zones. How do you configure the storage backend?

Answer:

1. Use a **multi-zone** storage class that allows for data replication across zones. For AWS, you can use **Amazon Elastic File System (EFS)** for multi-zone storage:

```
apiVersion: storage.k8s.io/v1
```

```
kind: StorageClass
```

```
metadata:
```

```
  name: efs-sc
```

```
provisioner: efs.csi.aws.com
```

```
parameters:
```

```
  provisioningMode: efs-ap
```

```
  filesystemId: fs-12345678
```

Kubernetes Networking

90. Scenario: You want to enforce strict east-west traffic control between microservices. How do you ensure traffic is routed securely and efficiently?

Answer:

1. **Service Mesh:** Deploy a service mesh like **Istio** to manage and secure microservice traffic.
 - Install Istio and enable sidecar injection for your namespaces:

```
istioctl install --set profile=default
```

```
kubectl label namespace <namespace> istio-injection=enabled
```

2. **Traffic Policies:** Use Istio's traffic management features to control how traffic flows between services, including setting up mutual TLS (mTLS) for encrypted communication.

```
apiVersion: security.istio.io/v1beta1
```

```
kind: PeerAuthentication
```

metadata:

name: default

namespace: default

spec:

mtls:

mode: STRICT

Kubernetes Autoscaling and Performance Optimization

91. Scenario: Your application experiences sudden spikes in traffic. How do you scale it automatically to handle increased load?

Answer:

1. **Horizontal Pod Autoscaler (HPA):** Set up HPA to scale the number of pod replicas based on CPU utilization or custom metrics:

```
kubectl autoscale deployment myapp --cpu-percent=50 --min=2 --max=10
```

2. **Cluster Autoscaler:** Ensure that your cluster can scale out nodes as needed by deploying the Cluster Autoscaler to add or remove nodes based on pod resource requests.

92. Scenario: Your cluster has a workload with unpredictable resource usage. How do you optimize resource allocation to avoid over-provisioning or under-provisioning?

Answer:

1. Use **Vertical Pod Autoscaler (VPA)** to adjust resource requests dynamically based on actual usage:
 - Deploy VPA for the workload:

```
apiVersion: autoscaling.k8s.io/v1
```

```
kind: VerticalPodAutoscaler
```

metadata:

name: myapp-vpa

spec:

targetRef:

apiVersion: "apps/v1"

kind: Deployment

name: myapp

updatePolicy:

updateMode: "Auto"

93. Scenario: You need to test the scalability and resilience of your application under load. How do you conduct stress testing?

Answer:

1. Use **load testing tools** like **Apache JMeter** or **K6** to simulate high traffic.
2. Create a dedicated load test namespace and deploy the tools there.
3. Monitor the cluster's response using **Prometheus** and **Grafana** to ensure autoscaling is triggered and that the application remains resilient under load.

Kubernetes Monitoring and Logging

94. Scenario: You need to collect and centralize application logs across your entire Kubernetes cluster. How do you configure a logging system?

Answer:

1. Deploy **Fluentd** or **Fluent Bit** as a DaemonSet on all nodes to collect logs.
2. Configure Fluentd to forward logs to **Elasticsearch** for storage and **Kibana** for visualization.
 - Example Fluentd configuration to send logs to Elasticsearch:

apiVersion: v1

kind: ConfigMap

metadata:

name: fluentd-config

data:

fluent.conf: |

<match **>

@type elasticsearch

host elasticsearch-logging

port 9200

</match>

95. Scenario: You need to monitor CPU, memory, and network usage for your Kubernetes cluster. How do you implement this?

Answer:

1. Install **Prometheus** to collect resource usage metrics.
2. Deploy **Grafana** to visualize these metrics using predefined dashboards for Kubernetes.
3. Install **Kube State Metrics** to gather detailed metrics about the state of Kubernetes resources.

Kubernetes CI/CD Integration

96. Scenario: You want to implement continuous deployment (CD) for your Kubernetes application with automatic rollbacks if the deployment fails. How do you configure this?

Answer:

1. **Jenkins Pipeline** or **GitLab CI**: Configure a CI/CD pipeline that automatically deploys the application on code changes.

- Example Jenkins pipeline:

```
pipeline {
  agent any
  stages {
    stage('Deploy') {
      steps {
        kubernetesDeploy(configs: 'deployment.yaml', kubeconfigId: 'kubeconfig')
      }
    }
  }
  post {
    failure {
      script {
        kubernetesRollback(kubeconfigId: 'kubeconfig', deployment: 'myapp')
      }
    }
  }
}
```

2. Use Kubernetes' **Deployment** strategy with a rolling update to ensure smooth transitions:

```
strategy:
  type: RollingUpdate
  rollingUpdate:
```

maxUnavailable: 1

maxSurge: 2

Kubernetes Disaster Recovery and Backups

97. Scenario: You need to back up all Kubernetes resources and restore them in case of disaster. How do you implement a backup and disaster recovery solution?

Answer:

1. Use **Velero** to back up the cluster's state (including resources and persistent volumes).
 - Install Velero with your cloud provider:

```
velero install --provider aws --bucket <BUCKET_NAME> --backup-location-config region=<REGION>
```

2. Create a backup:

```
velero backup create my-backup --include-namespaces mynamespace
```

3. Restore the backup:

```
velero restore create --from-backup my-backup
```

Kubernetes Security Best Practices

98. Scenario: Your organization has strict security requirements and must comply with CIS Kubernetes benchmarks. What are some key security practices you implement?

Answer:

1. **Control Plane Security:**
 - Enable role-based access control (RBAC) for the cluster.

- Ensure etcd data is encrypted in transit using TLS certificates.

2. Network Security:

- Implement **Network Policies** to restrict traffic between namespaces and pods.
- Deploy a **service mesh** like Istio to enforce mTLS for service-to-service communication.

3. Pod Security:

- Apply **PodSecurityPolicies (PSP)** to restrict the permissions and capabilities of running containers.
- Use **securityContext** in pods to run containers as non-root and apply read-only filesystem policies.

Kubernetes Security and Compliance

99. Scenario: Your Kubernetes cluster needs to comply with PCI-DSS for securing payment card information. What security features would you implement to meet compliance?

Answer:

1. **RBAC and Least Privilege:** Enforce **Role-Based Access Control (RBAC)** to ensure that users and service accounts have only the necessary permissions.
2. **PodSecurityPolicies (PSP) or Pod Security Standards (PSS):** Enforce pod security policies to limit the privileges of containers, such as:
 - Ensuring non-root containers.
 - Disabling privilege escalation.
 - Enforcing read-only file systems.
3. **Encryption:** Use **encryption at rest** for secrets and persistent volumes, ensuring that sensitive data like payment card information is encrypted.

4. **Network Segmentation:** Implement **Network Policies** to isolate payment services from the rest of the cluster.
5. **Audit Logging:** Enable Kubernetes audit logs to capture API requests and responses, which helps monitor access and detect any suspicious activity.

100. Scenario: You need to protect secrets in your Kubernetes cluster more securely than the default approach. What alternatives can you use?

Answer:

1. **HashiCorp Vault:** Use Vault to store and manage secrets outside of Kubernetes. Vault integrates with Kubernetes for secure secret management, allowing pods to dynamically request secrets.
2. **KMS (Key Management Service):** Use cloud provider-managed KMS (e.g., AWS KMS, Google Cloud KMS) to encrypt secrets before storing them in etcd.
 - For AWS, you can use KMS-integrated envelope encryption to encrypt Kubernetes secrets.
3. **Sealed Secrets:** Use Sealed Secrets by Bitnami to encrypt secrets in a GitOps workflow. The sealed secret can only be decrypted by the controller running in the cluster.

Kubernetes Ingress and Traffic Management

101. Scenario: Your application needs to handle SSL termination for multiple domains. How do you configure SSL certificates using Kubernetes Ingress?

Answer:

1. **Create Multiple TLS Secrets:** For each domain, generate or obtain an SSL certificate and create a Kubernetes Secret:

```
kubectl create secret tls example-com-tls --cert=example.com.crt --key=example.com.key
```

```
kubectl create secret tls example-org-tls --cert=example.org.crt --  
key=example.org.key
```

2. **Configure Ingress:** Define an Ingress resource that specifies both hosts and their corresponding TLS secrets:

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: multi-domain-ingress
```

```
spec:
```

```
  tls:
```

```
    - hosts:
```

```
      - example.com
```

```
      secretName: example-com-tls
```

```
    - hosts:
```

```
      - example.org
```

```
      secretName: example-org-tls
```

```
  rules:
```

```
    - host: example.com
```

```
      http:
```

```
        paths:
```

```
          - path: /
```

```
            pathType: Prefix
```

```
            backend:
```

```
              service:
```

```
                name: example-com-service
```

```
    port:
      number: 80
  - host: example.org
    http:
      paths:
      - path: /
        pathType: Prefix
      backend:
        service:
          name: example-org-service
          port:
            number: 80
```

102. Scenario: Your service needs to handle both HTTP and HTTPS traffic. How do you configure Ingress to support both?

Answer:

1. **Create TLS Secret:** Generate the SSL certificate and store it in a Kubernetes secret:

```
kubectl create secret tls my-tls-secret --cert=mydomain.com.crt --
key=mydomain.com.key
```

2. **Configure Ingress** with both HTTP and HTTPS paths:

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: http-https-ingress
```

```
spec:
```

tls:

- hosts:

- mydomain.com

secretName: my-tls-secret

rules:

- host: mydomain.com

http:

paths:

- path: /

pathType: Prefix

backend:

service:

name: myapp-service

port:

number: 80

103. Scenario: Your application has multiple versions, and you need to implement blue-green deployment with minimal downtime. How do you set this up using Kubernetes?

Answer:

1. **Create two separate deployments:** One for the current version (blue) and one for the new version (green).

- Example for the blue deployment:

apiVersion: apps/v1

kind: Deployment

metadata:

DevOps Shack

```
name: blue-deployment
```

```
spec:
```

```
replicas: 3
```

```
template:
```

```
  metadata:
```

```
    labels:
```

```
      app: myapp
```

```
      version: blue
```

```
  spec:
```

```
    containers:
```

```
      - name: myapp-container
```

```
        image: myapp:v1
```

- Similarly, create the green deployment with version v2.

2. **Use a Service:** Initially point the service to the blue deployment, then switch to the green deployment during the cutover:

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: myapp-service
```

```
spec:
```

```
  selector:
```

```
    app: myapp
```

```
    version: blue
```

During the cutover, update the selector to point to the green deployment:

spec:

selector:

app: myapp

version: green

Kubernetes Storage and Persistent Volumes

104. Scenario: You need to share data between pods in the same application. How do you set up shared storage?

Answer:

1. **Use a PersistentVolume (PV)** with an access mode of ReadWriteMany (RWX) to allow multiple pods to read and write to the same volume.

- Example PV configuration:

apiVersion: v1

kind: PersistentVolume

metadata:

name: shared-pv

spec:

capacity:

storage: 10Gi

accessModes:

- ReadWriteMany

nfs:

path: /exported/path

server: nfs-server.example.com

2. Create a **PersistentVolumeClaim (PVC)** that binds to the shared PV and use it in multiple pod specifications.
-

Kubernetes Monitoring and Logging

105. Scenario: You need to ensure that your Kubernetes cluster logs critical events like failed deployments and pod restarts. What monitoring and logging tools do you use?

Answer:

1. **Prometheus and Alertmanager:** Install Prometheus to collect metrics from the cluster and use Alertmanager to send notifications for critical events like pod crashes or deployment failures.

- Prometheus rule to detect pod restarts:

groups:

- name: pod-restart-rules

rules:

- alert: PodRestartAlert

expr: rate(kube_pod_container_status_restarts_total[5m]) > 0

for: 10m

labels:

severity: critical

annotations:

summary: "Pod has restarted"

2. **Fluentd or Fluent Bit:** Use Fluentd or Fluent Bit to collect logs from all pods and forward them to a centralized logging system like **Elasticsearch** or **Loki**.
3. **Grafana:** Use Grafana for visualizing metrics and logs with predefined Kubernetes dashboards for real-time monitoring.

Kubernetes Disaster Recovery

106. Scenario: Your cluster experiences a critical failure, and you need to restore it from backup. How do you perform disaster recovery?

Answer:

1. **Use Velero** to restore the cluster from backup:

- First, ensure that Velero is installed and configured to backup cluster resources.
- Create backups periodically:

```
velero backup create full-cluster-backup --include-namespaces '*'
```

2. **Restore from backup:**

```
velero restore create --from-backup full-cluster-backup
```

3. If persistent volumes are backed up, ensure they are restored as well. Velero handles both Kubernetes resource and PV backups.

Kubernetes Scaling and Performance

107. Scenario: Your application requires autoscaling based on custom metrics like request latency or queue length. How do you configure autoscaling with custom metrics?

Answer:

1. **Expose custom metrics:** Use a metrics exporter (e.g., Prometheus) to expose custom application metrics like queue length or request latency.
2. **Configure Horizontal Pod Autoscaler (HPA)** with custom metrics:

apiVersion: autoscaling/v2beta2

kind: HorizontalPodAutoscaler

metadata:

name: custom-hpa

spec:

scaleTargetRef:

apiVersion: apps/v1

kind: Deployment

name: myapp

minReplicas: 2

maxReplicas: 10

metrics:

- type: Pods

pods:

metric:

name: queue_length

target:

type: AverageValue

averageValue: "100"

3. Deploy **Prometheus Adapter** to expose custom metrics to the Kubernetes Metrics API for HPA to use.

Kubernetes Network and Traffic Management

108. Scenario: Your application requires strict network segmentation, and you need to control pod-to-pod communication based on namespaces and labels. How do you configure Network Policies?

Answer:

1. **NetworkPolicy:** Define a NetworkPolicy that allows ingress traffic only from pods in the same namespace with a specific label.

- Example NetworkPolicy:

apiVersion: networking.k8s.io/v1

kind: NetworkPolicy

metadata:

name: allow-same-namespace

namespace: app-namespace

spec:

podSelector:

matchLabels:

app: myapp

policyTypes:

- Ingress

ingress:

- from:

- podSelector:

matchLabels:

app: myapp

This ensures that only pods within the same namespace with the same label (app: myapp) can communicate.

Kubernetes CI/CD Integration

109. Scenario: You want to set up a CI/CD pipeline that builds, tests, and deploys a Kubernetes application automatically on every commit. How do you configure this?

Answer:

1. **Jenkins Pipeline:** Use Jenkins with Kubernetes plugin to deploy the application.

- Example Jenkins pipeline:

```
pipeline {
  agent any
  stages {
    stage('Build') {
      steps {
        sh 'docker build -t myapp:latest .'
        sh 'docker push myapp:latest'
      }
    }
    stage('Deploy') {
      steps {
        kubernetesDeploy(
          configs: 'k8s/deployment.yaml',
          kubeconfigId: 'kubeconfig'
```

```
)  
}  
}  
}  
}
```

2. Alternatively, use **GitLab CI** with built-in Kubernetes integration to automate builds and deployments.

Kubernetes Troubleshooting

110. Scenario: Your pods are getting evicted frequently due to memory pressure on the node. How do you resolve this?

Answer:

1. **Check Node Resource Usage:** Use `kubectl describe node <node-name>` to view memory and CPU pressure on the node.
2. **Set Resource Requests and Limits:** Ensure that pods have appropriate CPU and memory requests and limits configured to prevent over-allocation.
 - Example resource configuration:

resources:

requests:

memory: "100Mi"

cpu: "100m"

limits:

memory: "200Mi"

cpu: "200m"

3. **Vertical Pod Autoscaler (VPA):** Use VPA to automatically adjust pod resources based on actual usage.

111. Scenario: A pod is stuck in the Terminating state and won't delete. How do you troubleshoot and force delete it?

Answer:

1. **Check Pod Events:** Use `kubectl describe pod <pod-name>` to view events and identify why the pod is stuck.
2. **Force Delete** the pod if necessary:

```
kubectl delete pod <pod-name> --grace-period=0 --force
```

Kubernetes Performance Optimization

112. Scenario: You notice that certain nodes in your Kubernetes cluster are frequently overloaded with CPU usage. How do you resolve this issue?

Answer:

1. **Check Node Resource Usage:** Use `kubectl top nodes` to identify which nodes are experiencing high CPU usage.
2. **Redistribute Pods:** Use `kubectl drain <node-name>` to drain the overloaded node and let Kubernetes reschedule the pods to other nodes.
3. **Node Affinity/Anti-Affinity:** Configure pod `nodeAffinity` and `podAntiAffinity` to ensure balanced distribution of pods across nodes.

affinity:

podAntiAffinity:

requiredDuringSchedulingIgnoredDuringExecution:

- labelSelector:

matchExpressions:

- key: app

operator: In

values:

- myapp

topologyKey: "kubernetes.io/hostname"

4. **Autoscaling:** Enable **Cluster Autoscaler** to automatically add new nodes when the cluster is under heavy load.

113. Scenario: Your application is highly dependent on caching, and the pods are frequently being restarted, losing the cache data. How do you resolve this?

Answer:

1. Use **Persistent Volumes** to ensure cache data persists across pod restarts. Configure the pod to use a PVC (PersistentVolumeClaim):

volumeMounts:

- mountPath: /cache

name: cache-storage

volumes:

- name: cache-storage

persistentVolumeClaim:

claimName: cache-pvc

2. Consider using a **stateful application** like **Redis** or **Memcached** for caching, deployed as a StatefulSet to ensure data persistence.
3. Use **Memory Requests and Limits** to prevent the pod from being restarted due to resource constraints.

Kubernetes Load Balancing and Ingress

114. Scenario: You have multiple microservices running in your cluster, and you want to load balance traffic between them based on URL paths. How do you achieve this?

Answer:

1. Use **Ingress** to route traffic based on URL paths. Create an Ingress resource with path-based routing rules:

apiVersion: networking.k8s.io/v1

kind: Ingress

metadata:

name: myapp-ingress

spec:

rules:

- host: myapp.example.com

http:

paths:

- path: /service1

pathType: Prefix

backend:

service:

name: service1

port:

number: 80

- path: /service2

pathType: Prefix

backend:

service:

name: service2

port:

number: 80

2. Use an **Ingress Controller** such as **NGINX** to manage ingress traffic and enforce the rules.

115. Scenario: You need to expose an internal service to external users, but only allow access from a specific IP range. How do you configure this?

Answer:

1. Use a **LoadBalancer Service** for external access and configure **Network Policies** to allow access only from specific IP addresses.
2. Example LoadBalancer Service:

apiVersion: v1

kind: Service

metadata:

name: myapp-service

spec:

type: LoadBalancer

selector:

app: myapp

ports:

- protocol: TCP

port: 80

targetPort: 8080

3. Use a **NetworkPolicy** to restrict access:

apiVersion: networking.k8s.io/v1

kind: NetworkPolicy

metadata:

name: restrict-access

namespace: myapp

spec:

podSelector:

matchLabels:

app: myapp

ingress:

- from:

- ipBlock:

cidr: 192.168.1.0/24

Kubernetes Upgrades and Maintenance

116. Scenario: You want to upgrade your Kubernetes cluster to a new version with minimal downtime. How do you perform a safe upgrade?

Answer:

1. **Backup etcd:** Ensure you back up the etcd data before performing any upgrade:

etcdctl snapshot save snapshot.db

2. **Upgrade Control Plane Components:**

- Start by upgrading the API server, controller manager, and scheduler one node at a time:

kubeadm upgrade plan

kubeadm upgrade apply v1.x.x

3. Drain and Upgrade Worker Nodes:

- Drain nodes one by one and upgrade kubelet and kube-proxy:

kubectldrain <node-name> --ignore-daemonsets --force

apt-get update && apt-get install -y kubelet=1.x.x kube-proxy=1.x.x

kubectluncordon <node-name>

117. Scenario: You need to perform a rolling update of your application, but want to ensure that no more than 25% of the pods are unavailable during the update. How do you configure this?

Answer:

1. **Rolling Update Strategy:** Set the maxUnavailable and maxSurge parameters in your deployment strategy:

apiVersion: apps/v1

kind: Deployment

metadata:

name: myapp-deployment

spec:

replicas: 4

strategy:

type: RollingUpdate

rollingUpdate:

maxUnavailable: 25%

maxSurge: 1

This ensures that at least 75% of the pods are always available during the update.

Kubernetes Scaling

118. Scenario: You need to automatically scale the number of pods based on incoming HTTP requests. How do you configure autoscaling in Kubernetes?

Answer:

1. **Horizontal Pod Autoscaler (HPA):** Use HPA to scale pods based on CPU usage or custom metrics.

- Example HPA based on CPU utilization:

apiVersion: autoscaling/v1

kind: HorizontalPodAutoscaler

metadata:

name: myapp-hpa

spec:

scaleTargetRef:

apiVersion: apps/v1

kind: Deployment

name: myapp

minReplicas: 2

maxReplicas: 10

targetCPUUtilizationPercentage: 50

2. **Custom Metrics:** If you want to scale based on custom metrics such as request rate, install **Prometheus** to expose custom metrics and configure HPA with those metrics.

Kubernetes Storage Management

119. Scenario: Your application needs persistent storage that can be expanded dynamically. How do you configure this?

Answer:

1. Use a **StorageClass** that supports dynamic volume resizing, such as AWS EBS or GCP Persistent Disks.
2. Example StorageClass configuration for AWS EBS with volume expansion enabled:

```
apiVersion: storage.k8s.io/v1
```

```
kind: StorageClass
```

```
metadata:
```

```
  name: expandable-storage
```

```
provisioner: kubernetes.io/aws-ebs
```

```
parameters:
```

```
  type: gp2
```

```
  encrypted: "true"
```

```
allowVolumeExpansion: true
```

3. Modify the PersistentVolumeClaim (PVC) to request additional storage:

```
apiVersion: v1
```

```
kind: PersistentVolumeClaim
```

```
metadata:
```

```
  name: myapp-pvc
```

```
spec:
```

```
  storageClassName: expandable-storage
```

```
  accessModes:
```

- ReadWriteOnce

resources:

requests:

storage: 10Gi

120. Scenario: You need to replicate storage across multiple zones for high availability. How do you configure storage in Kubernetes for multi-zone replication?

Answer:

1. For AWS, use **Amazon EFS** or **EBS Multi-Attach** to allow multi-zone storage access.
2. Create a PersistentVolume with **ReadWriteMany (RWX)** access mode:

apiVersion: v1

kind: PersistentVolume

metadata:

name: efs-volume

spec:

capacity:

storage: 100Gi

accessModes:

- ReadWriteMany

nfs:

path: /exported/path

server: <efs-server>

Kubernetes CI/CD Pipelines

121. Scenario: You need to set up a CI/CD pipeline that deploys a Kubernetes application from GitHub automatically when new commits are pushed. How do you configure this?

Answer:

1. **Jenkins Integration:** Set up a Jenkins pipeline that triggers on GitHub push events.

- Example Jenkinsfile for building and deploying the application:

```
pipeline {
  agent any
  stages {
    stage('Build') {
      steps {
        sh 'docker build -t myapp:latest .'
        sh 'docker push myapp:latest'
      }
    }
    stage('Deploy') {
      steps {
        kubernetesDeploy(
          configs: 'k8s/deployment.yaml',
          kubeconfigId: 'kubeconfig'
        )
      }
    }
  }
}
```

```
}  
}
```

2. Alternatively, use **GitLab CI** with Kubernetes integration to automate deployment on code changes.

Kubernetes Security

122. Scenario: You want to ensure that no container runs with root privileges in your Kubernetes cluster. How do you enforce this?

Answer:

1. **PodSecurityPolicies (PSP):** Use PodSecurityPolicies or Pod Security Standards to enforce non-root containers.
 - Example PSP configuration:

```
apiVersion: policy/v1beta1
```

```
kind: PodSecurityPolicy
```

```
metadata:
```

```
  name: non-root-policy
```

```
spec:
```

```
  privileged: false
```

```
  runAsUser:
```

```
    rule: MustRunAsNonRoot
```

```
  fsGroup:
```

```
    rule: MustRunAs
```

2. Bind the PSP to roles or service accounts using **RoleBinding** or **ClusterRoleBinding**.

123. Scenario: You want to restrict container capabilities and ensure that containers cannot escalate privileges. How do you configure this?

Answer:

1. Use the securityContext to restrict container capabilities and prevent privilege escalation:

securityContext:

runAsNonRoot: true

capabilities:

drop:

- ALL

allowPrivilegeEscalation: false

2. Apply this securityContext to each container in the pod specification to enforce security best practices.

Kubernetes Disaster Recovery

124. Scenario: Your etcd cluster has become corrupted, and you need to restore it from a backup. What are the steps for disaster recovery?

Answer:

1. **Restore etcd from backup:**

- Stop the etcd service and restore from the snapshot:

```
etcdctl snapshot restore snapshot.db --data-dir=/var/lib/etcd
```

2. **Restart etcd:**

- After restoring the backup, restart etcd and ensure that the cluster is functional.

3. Verify the health of etcd and check the cluster state:

```
etcdctl endpoint health
```

```
kubectl get nodes
```

125. Scenario: You need to perform full-cluster backups and restore critical Kubernetes resources in case of a failure. What tool would you use, and how do you configure it?

Answer:

1. Use **Velero** to perform cluster backups and restore operations.

- Install Velero:

```
velero install --provider aws --bucket <bucket-name> --backup-location-config region=<region>
```

2. Perform a full-cluster backup:

```
velero backup create cluster-backup --include-namespaces '*'
```

3. To restore the backup:

```
velero restore create --from-backup cluster-backup
```

Kubernetes Multi-Cluster Management

126. Scenario: Your organization runs multiple Kubernetes clusters across different regions and cloud providers. How do you manage these clusters centrally?

Answer:

1. **KubeFed (Kubernetes Federation):** Use **KubeFed** to centrally manage multiple clusters, ensuring consistent configurations and deployments across regions.

- Install KubeFed and join clusters:

```
kubefedctl join <cluster-name> --host-cluster-context <host-cluster>
```

2. **Helm Charts:** Package applications as Helm charts to enable consistent deployment across multiple clusters.

3. **GitOps Tools:** Use GitOps tools like **ArgoCD** or **Flux** to manage multiple clusters declaratively, ensuring a single source of truth (via Git) for application deployments.

127. Scenario: You need to deploy an application in a multi-region setup with traffic distributed across clusters based on user location. How do you achieve this?

Answer:

1. Use **Global Load Balancers** provided by cloud providers (e.g., AWS Global Accelerator, Google Cloud Global Load Balancer) to route traffic to the nearest region.
2. **KubeFed:** Federate the clusters across regions and deploy the application using **Kubernetes Federation (KubeFed)** to manage consistent configurations.
3. For **Istio Service Mesh** users:
 - Use **Istio** to manage multi-cluster deployments, enabling cross-cluster service discovery and load balancing between clusters.

Kubernetes Security Best Practices

128. Scenario: You need to secure sensitive environment variables (e.g., API keys, database credentials) used by pods in Kubernetes. How do you achieve this?

Answer:

1. **Kubernetes Secrets:** Store sensitive data in **Kubernetes Secrets** and mount them as environment variables or volumes in pods.
 - Example Secret:

DevOps Shack

apiVersion: v1

kind: Secret

metadata:

name: api-credentials

type: Opaque

data:

api-key: YXBpLWtleQ== # base64 encoded value

2. **Service Mesh with mTLS:** If using a service mesh like **Istio**, enable **mutual TLS (mTLS)** to encrypt communication between services.
3. **Vault Integration:** Use **HashiCorp Vault** or cloud-based key management services (AWS KMS, GCP KMS) for storing secrets and dynamically injecting them into pods at runtime.

129. Scenario: Your organization requires compliance with security standards, and you need to audit API access to the Kubernetes cluster. How do you enable audit logging?

Answer:

1. **Enable Kubernetes Audit Logs:** Configure Kubernetes to capture audit logs of API server requests.
 - Example audit policy file:

apiVersion: audit.k8s.io/v1

kind: Policy

rules:

- level: RequestResponse

resources:

- group: ""

resources: ["pods", "secrets", "configmaps"]

DevOps Shack

- level: Metadata

resources:

- group: ""

resources: ["services", "namespaces"]

Apply the audit policy by passing the file to the API server:

bash

Copy code

--audit-policy-file=/etc/kubernetes/audit-policy.yaml

2. Use **centralized logging solutions** (e.g., Elasticsearch, Fluentd) to forward audit logs to a secure storage location for analysis and auditing.

130. Scenario: You want to restrict certain users to specific namespaces and allow only read access to Kubernetes resources. How do you implement this using RBAC?

Answer:

1. **Create a Role** with read-only permissions:

kind: Role

apiVersion: rbac.authorization.k8s.io/v1

metadata:

namespace: dev

name: read-only-role

rules:

- apiGroups: [""]

resources: ["pods", "services"]

verbs: ["get", "list", "watch"]

2. **Bind the Role** to the specific user:

bash

Copy code

```
kubectl create rolebinding read-only-binding --role=read-only-role --  
user=username --namespace=dev
```

131. Scenario: You need to implement a policy that ensures all containers in your cluster run with a read-only filesystem. How do you enforce this?

Answer:

1. **Pod Security Policies (PSP):** Create a PodSecurityPolicy that enforces a read-only filesystem for all containers:

```
apiVersion: policy/v1beta1
```

```
kind: PodSecurityPolicy
```

```
metadata:
```

```
  name: enforce-read-only
```

```
spec:
```

```
  requiredDropCapabilities:
```

```
  - ALL
```

```
  readOnlyRootFilesystem: true
```

2. Apply the policy to relevant namespaces by binding the PSP to roles and users using **RoleBindings** or **ClusterRoleBindings**.

Kubernetes Logging and Monitoring

132. Scenario: You need to collect logs from all containers and centralize them in a searchable interface for analysis. What logging solution do you configure?

Answer:

1. Deploy **Fluentd** or **Fluent Bit** as a **DaemonSet** to collect logs from all nodes in the Kubernetes cluster.

- Example DaemonSet for Fluentd:

apiVersion: apps/v1

kind: DaemonSet

metadata:

name: fluentd

namespace: logging

spec:

selector:

matchLabels:

name: fluentd

template:

metadata:

labels:

name: fluentd

spec:

containers:

- name: fluentd

image: fluent/fluentd:v1.11

env:

- name: FLUENT_ELASTICSEARCH_HOST

value: "elasticsearch-logging"

2. Send logs to **Elasticsearch** and use **Kibana** for visualization and log searching.
3. Use **Grafana Loki** as an alternative for lightweight logging, integrating logs with metrics.

Kubernetes Performance Tuning

133. Scenario: Your application is facing performance issues due to high memory usage, causing pods to be killed frequently. How do you address this?

Answer:

1. **Resource Requests and Limits:** Ensure that pods have proper memory and CPU requests and limits defined to avoid resource over-allocation:

resources:

requests:

memory: "512Mi"

cpu: "500m"

limits:

memory: "1024Mi"

cpu: "1000m"

2. Use **Vertical Pod Autoscaler (VPA)** to dynamically adjust the resource limits based on actual usage patterns.

- Install VPA:

kubectl apply -f

<https://github.com/kubernetes/autoscaler/releases/download/v0.9.0/vertical-pod-autoscaler.yaml>

- Create VPA resource:

apiVersion: autoscaling.k8s.io/v1

kind: VerticalPodAutoscaler

metadata:

name: myapp-vpa

spec:

targetRef:

apiVersion: "apps/v1"

kind: Deployment

name: myapp

updatePolicy:

updateMode: "Auto"

134. Scenario: Your application requires faster scaling in response to sudden traffic spikes. How do you ensure Kubernetes autoscaling responds quickly?

Answer:

1. **Tune HPA (Horizontal Pod Autoscaler)** scaling policies:

- Reduce the **stabilizationWindow** and increase the **syncPeriod** to allow quicker responses to scaling events.

apiVersion: autoscaling/v2beta2

kind: HorizontalPodAutoscaler

metadata:

name: myapp-hpa

spec:

behavior:

scaleUp:

stabilizationWindowSeconds: 30

policies:

- type: Percent

value: 100

periodSeconds: 60

minReplicas: 2

maxReplicas: 20

metrics:

- type: Resource

resource:

name: cpu

target:

type: Utilization

averageUtilization: 50

Kubernetes CI/CD Pipelines

135. Scenario: You need to implement a CI/CD pipeline that automatically builds and deploys your Kubernetes application after each commit, ensuring quality through unit tests and security checks. How do you set this up?

Answer:

1. Jenkins Integration:

- Set up a Jenkins pipeline with the following stages:
 - **Build Stage:** Build the Docker image and push it to a container registry.
 - **Test Stage:** Run unit tests and vulnerability scans (e.g., with Trivy).
 - **Deploy Stage:** Deploy the updated image to the Kubernetes cluster.
- Example Jenkinsfile:

```
pipeline {  
  agent any  
  stages {
```

```
stage('Build') {
  steps {
    sh 'docker build -t myapp:latest .'
    sh 'docker push myapp:latest'
  }
}
stage('Test') {
  steps {
    sh 'run-unit-tests.sh'
    sh 'trivy image --exit-code 1 myapp:latest'
  }
}
stage('Deploy') {
  steps {
    kubernetesDeploy(
      configs: 'k8s/deployment.yaml',
      kubeconfigId: 'kubeconfig'
    )
  }
}
}
```

2. **GitLab CI/CD:** Use GitLab CI/CD pipelines for similar functionality with native Kubernetes integration.

Kubernetes High Availability (HA)

136. Scenario: You need to ensure high availability (HA) for the Kubernetes control plane and ensure that the etcd cluster is fault-tolerant. How do you set this up?

Answer:

1. **Multiple Master Nodes:** Set up multiple master nodes (API servers, controller managers, and schedulers) across different availability zones or data centers. Use a load balancer in front of the API servers.
2. **HA etcd Cluster:** Set up a highly available etcd cluster with an odd number of members (e.g., 3 or 5) to ensure quorum-based consensus.

- Example etcd cluster configuration:

```
etcd --name etcd1 --initial-advertise-peer-urls http://etcd1.example.com:2380 \
--listen-peer-urls http://0.0.0.0:2380 \
--initial-cluster
etcd1=http://etcd1.example.com:2380,etcd2=http://etcd2.example.com:2380,et
cd3=http://etcd3.example.com:2380 \
--initial-cluster-token etcd-cluster-1 \
--initial-cluster-state new
```

137. Scenario: Your Kubernetes worker nodes experience periodic failures, leading to pod evictions. How do you ensure that workloads remain available during node failures?

Answer:

1. **Pod Disruption Budgets (PDBs):** Use **Pod Disruption Budgets** to ensure that a minimum number of replicas remain available during node failures or maintenance:

apiVersion: policy/v1

kind: PodDisruptionBudget

metadata:

name: myapp-pdb

spec:

minAvailable: 80%

selector:

matchLabels:

app: myapp

2. **Cluster Autoscaler:** Use Cluster Autoscaler to automatically replace failed nodes and reschedule pods.
3. **Node Pools:** Configure **node pools** with redundant capacity to handle workloads during node failures.

Kubernetes Network and Traffic Management

138. Scenario: You need to ensure secure traffic between microservices in your Kubernetes cluster. How do you configure secure communication between services?

Answer:

1. **Service Mesh with mTLS:** Use **Istio** or **Linkerd** to enforce **mutual TLS (mTLS)** between microservices.
 - Example Istio configuration for mTLS:

apiVersion: security.istio.io/v1beta1

kind: PeerAuthentication

metadata:

name: default

namespace: default

spec:

mtls:

mode: STRICT

2. **Network Policies:** Use **Network Policies** to restrict communication between services to only authorized pods or namespaces.

139. Scenario: You need to enforce rate limiting on specific services to prevent abuse. How do you implement this in Kubernetes?

Answer:

1. **Service Mesh:** Use **Istio** to configure rate limiting for specific services.
 - Define a QuotaSpec to limit the number of requests:

apiVersion: config.istio.io/v1alpha2

kind: QuotaSpec

metadata:

name: request-count

namespace: default

spec:

rules:

- quotas:

- quota: requestcount

charge: 1

- Use QuotaSpecBinding to apply it to services.
2. **Ingress Controller:** Configure rate limiting using **NGINX Ingress Controller**:
 - Example annotations:

metadata:

annotations:

nginx.ingress.kubernetes.io/limit-connections: "20"

nginx.ingress.kubernetes.io/limit-rps: "10"

140. Scenario: Your services need to communicate across different Kubernetes clusters in different regions. How do you handle cross-cluster service communication?

Answer:

1. **Istio Multi-Cluster:** Use **Istio** to create a multi-cluster service mesh. Istio enables cross-cluster communication by creating a shared service registry and implementing secure communication between services in different clusters.
 - Configure Istio control planes in each cluster and set up a shared root CA for mTLS.
2. **KubeFed:** Alternatively, use **KubeFed** for cross-cluster service discovery and load balancing.

Kubernetes Stateful Workloads and Databases

141. Scenario: You need to deploy a highly available stateful database, such as PostgreSQL, in Kubernetes. How do you ensure data persistence and fault tolerance?

Answer:

1. **StatefulSets:** Use **StatefulSets** to deploy PostgreSQL with persistent storage and stable network identities.
 - Example PostgreSQL StatefulSet:

DevOps Shack

apiVersion: apps/v1

kind: StatefulSet

metadata:

name: postgres

spec:

serviceName: "postgres"

replicas: 3

selector:

matchLabels:

app: postgres

template:

metadata:

labels:

app: postgres

spec:

containers:

- name: postgres

image: postgres:12

volumeMounts:

- name: pgdata

mountPath: /var/lib/postgresql/data

volumeClaimTemplates:

- metadata:

name: pgdata

spec:

accessModes: ["ReadWriteOnce"]

resources:

requests:

storage: 10Gi

2. **Persistent Volumes:** Ensure persistent volumes (PVs) are created and bound to ensure data persistence across pod restarts.

142. Scenario: Your database needs to be scaled vertically (increased CPU and memory) during periods of high load. How do you manage this in Kubernetes?

Answer:

1. **Vertical Pod Autoscaler (VPA):** Use VPA to automatically adjust CPU and memory requests/limits for database pods:

- Install VPA and create a VPA resource:

apiVersion: autoscaling.k8s.io/v1

kind: VerticalPodAutoscaler

metadata:

name: postgres-vpa

spec:

targetRef:

apiVersion: apps/v1

kind: StatefulSet

name: postgres

updatePolicy:

updateMode: "Auto"

Kubernetes Backup and Disaster Recovery

143. Scenario: You need to implement a backup solution that ensures both Kubernetes resources and persistent volumes can be restored in case of disaster. What tool would you use and how?

Answer:

1. **Velero:** Use **Velero** to back up both Kubernetes resources and persistent volumes.

- Install Velero with your cloud provider (AWS, GCP, Azure):

```
velero install --provider aws --bucket <bucket-name> --backup-location-config region=<region>
```

- Backup the entire cluster:

```
velero backup create full-cluster-backup --include-namespaces '*'
```

- To restore:

```
velero restore create --from-backup full-cluster-backup
```

144. Scenario: You need to schedule automated backups for your Kubernetes resources daily. How do you configure this with Velero?

Answer:

1. **Schedule Backups:** Use Velero to create a scheduled backup job.

```
velero create schedule daily-backup --schedule "0 2 * * *" --include-namespaces '*'
```

Kubernetes Security and Compliance

145. Scenario: Your organization needs to enforce network security policies that restrict pod communication across namespaces. How do you configure this?

Answer:

1. **Network Policies:** Use **Kubernetes Network Policies** to enforce network isolation across namespaces.

- Example Network Policy to restrict ingress from other namespaces:

apiVersion: networking.k8s.io/v1

kind: NetworkPolicy

metadata:

name: restrict-namespace-access

spec:

podSelector: {}

policyTypes:

- Ingress

ingress:

- from:

- namespaceSelector:

- matchLabels:

- name: mynamespace

146. Scenario: Your cluster needs to comply with HIPAA security standards. How do you ensure secure handling of sensitive data in Kubernetes?

Answer:

1. **Kubernetes Secrets:** Store sensitive data, such as credentials, in **Kubernetes Secrets**.
2. **Encryption at Rest:** Enable **encryption at rest** for etcd to protect sensitive data stored in Kubernetes Secrets.
3. **Audit Logs:** Enable **Kubernetes audit logging** to track API access and monitor any potential breaches of sensitive data.

Kubernetes Networking Troubleshooting

147. Scenario: Pods in your cluster are unable to communicate with external services. How do you troubleshoot this issue?

Answer:

1. **Check Network Policies:** Ensure that there are no Network Policies blocking egress traffic.
2. **DNS Resolution:** Verify that DNS is working within the cluster by running:
`kubect exec -it <pod> -- nslookup google.com`
3. **Check Service Configuration:** Ensure the ClusterIP or LoadBalancer service is correctly configured to expose the pod.
4. **kube-proxy:** Ensure that **kube-proxy** is running and configured correctly to handle network traffic between pods and external services.

148. Scenario: You notice that DNS lookups are failing intermittently in your cluster. How do you diagnose and resolve the issue?

Answer:

1. **Check CoreDNS Pods:** Ensure that **CoreDNS** pods are running and healthy:

`kubect exec -it <pod> -- nslookup google.com`

2. **CoreDNS ConfigMap:** Check the CoreDNS configuration in the kube-system namespace for any misconfigurations:

`kubect exec -it <pod> -- nslookup google.com`

Kubernetes Cluster Management and Scaling

149. Scenario: You need to scale your Kubernetes cluster nodes automatically based on workload demand. How do you implement this?

Answer:

1. **Cluster Autoscaler:** Install and configure **Cluster Autoscaler** to add or remove nodes based on resource utilization.

- For GKE:

```
gcloud container clusters update my-cluster --enable-autoscaling --min-nodes=3 --max-nodes=10
```

2. **Ensure Resource Requests:** Ensure that pods have CPU and memory requests/limits defined so that the autoscaler can accurately assess node usage.

150. Scenario: You notice that your cluster is not scaling out even when CPU utilization is high. What could be the issue, and how do you troubleshoot it?

Answer:

1. **Check HPA Configuration:** Ensure that the **Horizontal Pod Autoscaler (HPA)** is configured correctly and targeting the right metrics:

```
kubectl get hpa
```

2. **Cluster Autoscaler Logs:** Check the Cluster Autoscaler logs to ensure it is functioning correctly and can provision new nodes.
3. **Resource Requests:** Ensure pods have CPU requests defined. If no CPU requests are defined, HPA cannot trigger scaling actions.

Kubernetes Performance Troubleshooting

151. Scenario: Your application experiences slow response times during peak load, and some requests are being dropped. How do you troubleshoot and resolve this?

Answer:

1. **Resource Limits:** Check if the pods have hit their resource limits:

```
kubectl describe pod <pod-name>
```

2. **Horizontal Pod Autoscaler:** Ensure that HPA is scaling the pods based on CPU or custom metrics like request rate.
3. **Network Latency:** Check for network bottlenecks by inspecting pod-to-pod latency. Use tools like **Weave Scope** or **Prometheus** to visualize network traffic.
4. **Pod Logs:** Check logs for any application-level issues such as timeouts or connection errors.

152. Scenario: Some of your nodes are running out of disk space, causing pods to be evicted. How do you manage disk usage in Kubernetes?

Answer:

1. **Log Rotation:** Implement log rotation in the container runtime (e.g., Docker, containerd) to prevent large log files from filling the disk.
2. **Monitor Disk Usage:** Use **Prometheus** to monitor disk usage and set up alerts when nodes approach capacity.
3. **Eviction Policies:** Review your **Kubelet Eviction Policies** to ensure that pods are not evicted prematurely due to disk pressure.

Kubernetes Troubleshooting - Pod and Container Issues

153. Scenario: A pod is stuck in Pending state due to insufficient CPU or memory. How do you resolve this?

Answer:

1. **Check Events:** Run `kubectl describe pod <pod-name>` to check events related to resource scheduling failures.
2. **Node Resource Availability:** Verify if nodes have enough resources by checking `kubectl top nodes`.
3. **Horizontal Pod Autoscaler:** If HPA is not scaling the application, check the HPA configuration and ensure resource requests are defined in the deployment spec.

154. Scenario: Your pods are being frequently evicted due to high memory usage. How do you identify the issue and prevent future evictions?

Answer:

1. **Check Pod Events:** Use `kubectl describe pod <pod-name>` to view the pod's eviction events.
2. **Memory Limits:** Set appropriate memory requests and limits to prevent pods from overconsuming memory:

resources:

requests:

memory: "512Mi"

limits:

memory: "1024Mi"

3. **Vertical Pod Autoscaler (VPA):** Use VPA to dynamically adjust memory requests based on the pod's usage patterns.

Kubernetes Observability and Monitoring

155. Scenario: You need to monitor application performance and set up alerts for key metrics like CPU utilization and request latency. How do you configure monitoring?

Answer:

1. **Prometheus:** Set up **Prometheus** to scrape metrics from Kubernetes and application components.
 - Use `kube-state-metrics` for cluster-level metrics.
2. **Grafana Dashboards:** Set up **Grafana** with pre-built Kubernetes dashboards to visualize cluster and application metrics.
3. **Alertmanager:** Use **Alertmanager** to set up alerts for critical metrics:

- Example CPU utilization alert:

groups:

- name: cpu-usage-alert

rules:

- alert: HighCPUUsage

expr: 100 - (avg by (instance)
(rate(node_cpu_seconds_total{mode="idle"}[5m])) * 100) > 80

for: 10m

labels:

severity: critical

annotations:

summary: "Instance {{ \$labels.instance }} CPU usage is high"

Kubernetes Ingress and Load Balancing

156. Scenario: You want to configure TLS termination for multiple services behind a single Ingress controller. How do you achieve this?

Answer:

1. **Create TLS Secrets:** Create TLS secrets for each domain or service that requires TLS termination:

```
kubectl create secret tls tls-secret --cert=tls.crt --key=tls.key
```

2. **Configure Ingress Resource:** Set up an Ingress resource that handles multiple services and includes TLS termination for each domain:

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
name: myapp-ingress
spec:
  tls:
  - hosts:
    - myapp.com
    secretName: tls-secret
  rules:
  - host: myapp.com
    http:
      paths:
      - path: /
        backend:
          service:
            name: myapp-service
            port:
              number: 80
```

157. Scenario: You want to implement a canary deployment to gradually shift traffic from an older version of your service to a new version. How do you configure this with Kubernetes?

Answer:

1. **Service with Weighted Traffic Splitting:** Use **Istio** or **NGINX Ingress Controller** to implement traffic splitting between the two versions.
 - Example Istio VirtualService for canary deployment:

```
apiVersion: networking.istio.io/v1alpha3
```

```
kind: VirtualService
```

metadata:

name: myapp

spec:

hosts:

- myapp.example.com

http:

- route:

- destination:

host: myapp

subset: v1

weight: 90

- destination:

host: myapp

subset: v2

weight: 10

Kubernetes Cluster Upgrades

158. Scenario: You want to upgrade your Kubernetes cluster to a newer version with minimal downtime. What are the steps you follow?

Answer:

1. **Backup etcd:** Ensure you have an etcd backup before upgrading the control plane.
2. **Upgrade the Control Plane:** Use **kubeadm** to upgrade the control plane components (API server, controller manager, and scheduler).

```
kubeadm upgrade apply v1.x.x
```

3. **Upgrade Worker Nodes:** Drain and upgrade worker nodes one at a time:

```
kubectldrain <node-name> --ignore-daemonsets --force
```

```
apt-get update && apt-get install -y kubelet=1.x.x kubeadm=1.x.x
```

```
kubectluncordon <node-name>
```

Kubernetes Storage and Persistent Volumes

159. Scenario: You need to dynamically provision persistent volumes for your application using cloud storage (AWS EBS, GCP PD). How do you configure this?

Answer:

1. **StorageClass:** Define a StorageClass that provisions storage using the cloud provider's storage service (e.g., AWS EBS or GCP PD).

- Example StorageClass for AWS EBS:

```
apiVersion: storage.k8s.io/v1
```

```
kind: StorageClass
```

```
metadata:
```

```
  name: ebs-storage
```

```
provisioner: kubernetes.io/aws-ebs
```

```
parameters:
```

```
  type: gp2
```

```
  fsType: ext4
```

160. Scenario: Your application requires shared storage between multiple pods. How do you configure shared persistent storage in Kubernetes?

Answer:

1. **PersistentVolume with RWX:** Create a PersistentVolume with ReadWriteMany (RWX) access mode.

- Example NFS-backed PersistentVolume:

apiVersion: v1

kind: PersistentVolume

metadata:

name: shared-volume

spec:

capacity:

storage: 10Gi

accessModes:

- ReadWriteMany

nfs:

path: /exported/path

server: <nfs-server>

Kubernetes Custom Metrics and Autoscaling

161. Scenario: You want to scale your application based on custom metrics such as request count or latency. How do you implement this?

Answer:

1. **Prometheus Adapter:** Install **Prometheus Adapter** to expose custom metrics to the Kubernetes Metrics API.
 - Install the Prometheus Adapter:

```
kubectl apply -f https://github.com/DirectXMan12/k8s-prometheus-adapter/releases/download/v0.5.0/prometheus-adapter.yaml
```

2. **Horizontal Pod Autoscaler (HPA):** Create an HPA that scales based on the custom metrics:

```
apiVersion: autoscaling/v2beta2
```

```
kind: HorizontalPodAutoscaler
```

```
metadata:
```

```
  name: myapp-hpa
```

```
spec:
```

```
  minReplicas: 2
```

```
  maxReplicas: 10
```

```
  metrics:
```

```
    - type: Pods
```

```
      pods:
```

```
        metric:
```

```
          name: request_count
```

```
        target:
```

```
          type: AverageValue
```

```
          averageValue: "100"
```

Kubernetes Service Mesh and Traffic Management

162. Scenario: You need to ensure secure, encrypted communication between microservices in your cluster. How do you achieve this with Istio?

Answer:

1. **Install Istio:** Set up Istio as a service mesh to manage secure communication between microservices.
2. **Mutual TLS (mTLS):** Enable mTLS to encrypt traffic between services.
 - Example PeerAuthentication for mTLS:

apiVersion: security.istio.io/v1beta1

kind: PeerAuthentication

metadata:

name: default

namespace: default

spec:

mtls:

mode: STRICT

163. Scenario: You need to perform advanced traffic shaping, such as fault injection and retries, between services. How do you implement this in Istio?

Answer:

1. **Istio VirtualService:** Use **Istio VirtualService** to configure traffic management policies such as retries and fault injection.
 - Example VirtualService with retries and fault injection:

apiVersion: networking.istio.io/v1alpha3

kind: VirtualService

metadata:

name: myapp

spec:

hosts:

- myapp.example.com

http:

- route:

- destination:

host: myapp

retries:

attempts: 3

perTryTimeout: 2s

fault:

abort:

httpStatus: 500

percentage:

value: 10

Kubernetes Advanced Topics

164. Scenario: Your application needs to handle batch processing workloads that can be scheduled at specific times. How do you configure this in Kubernetes?

Answer:

1. **CronJob:** Use **Kubernetes CronJob** to schedule periodic batch jobs.

○ Example CronJob:

apiVersion: batch/v1

kind: CronJob

metadata:

name: batch-job

DevOps Shack

spec:

schedule: "0 2 * * *"

jobTemplate:

spec:

template:

spec:

containers:

- name: batch-container

image: mybatchimage

command: ["run-batch-job.sh"]

restartPolicy: OnFailure

165. Scenario: Your application includes GPU workloads for machine learning. How do you configure GPU support in your Kubernetes cluster?

Answer:

1. **Install GPU Drivers:** Install GPU drivers on your Kubernetes nodes (e.g., NVIDIA drivers for NVIDIA GPUs).
2. **Enable GPU Scheduling:** Use the `nvidia.com/gpu` resource in your pod specification to schedule GPU workloads.
 - Example pod with GPU:

apiVersion: v1

kind: Pod

metadata:

name: gpu-pod

spec:

containers:

```
- name: gpu-container  
  image: nvidia/cuda:10.0  
  resources:  
    limits:  
      nvidia.com/gpu: 1
```

Kubernetes Disaster Recovery

166. Scenario: Your etcd database becomes corrupted, and you need to restore it from a backup. What are the steps to restore etcd in Kubernetes?

Answer:

1. **Stop etcd:** Stop the etcd service on the master nodes.
2. **Restore from Snapshot:**
 - Use etcdctl to restore etcd from a backup:

```
etcdctl snapshot restore snapshot.db --data-dir /var/lib/etcd
```

3. **Restart etcd:** Restart etcd and verify the health of the cluster.

167. Scenario: You need to ensure that your entire Kubernetes cluster can be restored in case of disaster. What tools and steps would you use?

Answer:

1. **Velero:** Use **Velero** to back up and restore both Kubernetes resources and persistent volumes.
 2. **Automate Backups:** Set up scheduled backups with Velero and ensure backup storage is replicated across regions or availability zones.
-

Kubernetes Cluster Federation

168. Scenario: You need to deploy your application across multiple Kubernetes clusters in different regions with unified management. How do you implement this?

Answer:

1. **KubeFed (Kubernetes Federation):** Use **KubeFed** to manage multiple clusters from a single control plane.

- Install KubeFed:

```
kubectl apply -f https://github.com/kubernetes-sigs/kubefed/releases/download/v0.1.0/kubefed.yaml
```

- Join clusters to the federation:

```
kubefedctl join <cluster-name> --host-cluster-context <host-cluster>
```

2. **Multi-Cluster Ingress:** Use **multi-cluster ingress controllers** to manage ingress traffic across multiple regions.

Kubernetes Role-Based Access Control (RBAC)

169. Scenario: You want to limit a team's access to only manage deployments in a specific namespace without giving access to other namespaces. How do you configure RBAC for this?

Answer:

1. **Create Role:** Define a role with limited permissions to manage deployments in the specific namespace.

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: Role
```

```
metadata:
```

```
  namespace: dev
```

```
name: deployment-manager
rules:
- apiGroups: ["apps"]
  resources: ["deployments"]
  verbs: ["create", "update", "delete", "get"]
```

2. **RoleBinding:** Bind the role to specific users or groups using a RoleBinding.

```
kubectl create rolebinding deployment-access --role=deployment-manager --
user=team-member --namespace=dev
```

170. Scenario: You want to allow a service account to create and delete ConfigMaps in a specific namespace but not other resources. How do you configure this?

Answer:

1. **Create Role:** Define a role with permissions to manage ConfigMaps only:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: dev
  name: configmap-manager
rules:
- apiGroups: [""]
  resources: ["configmaps"]
  verbs: ["create", "delete", "get", "list"]
```

2. **RoleBinding:** Bind the role to a service account:

```
kubectl create rolebinding configmap-access --role=configmap-manager --
serviceaccount=dev:my-service-account --namespace=dev
```

Kubernetes Troubleshooting Advanced

171. Scenario: Your application is experiencing intermittent crashes and the logs indicate out-of-memory (OOM) errors. How do you diagnose and resolve this issue?

Answer:

1. **Check Pod Events:** Use `kubectl describe pod <pod-name>` to check for OOM kill events.
2. **Set Memory Limits:** Ensure memory limits are set for your application to prevent uncontrolled memory usage:

resources:

requests:

memory: "512Mi"

limits:

memory: "1024Mi"

3. **Monitor Memory Usage:** Use **Prometheus** to monitor memory usage and alert on high usage levels.

172. Scenario: A pod is stuck in ContainerCreating due to issues with PersistentVolume mounts. How do you troubleshoot and resolve this?

Answer:

1. **Check Pod Events:** Use `kubectl describe pod <pod-name>` to view events related to volume mounting.
2. **PersistentVolumeClaim:** Ensure that the PersistentVolumeClaim (PVC) is correctly bound to a PersistentVolume (PV):

```
kubectl get pvc <pvc-name>
```

Kubernetes Networking - Advanced

173. Scenario: You want to allow only specific external IP addresses to access your application via Ingress. How do you configure this?

Answer:

1. **Whitelist IPs in Ingress Annotations:** Use **NGINX Ingress Controller** annotations to whitelist specific IP addresses:

```
apiVersion: networking.k8s.io/v1
```

```
kind: Ingress
```

```
metadata:
```

```
  name: myapp-ingress
```

```
  annotations:
```

```
    nginx.ingress.kubernetes.io/whitelist-source-range: "192.168.1.0/24"
```

```
spec:
```

```
  rules:
```

```
  - host: myapp.example.com
```

```
    http:
```

```
      paths:
```

```
      - path: /
```

```
        backend:
```

```
          service:
```

```
            name: myapp-service
```

```
            port:
```

```
              number: 80
```

174. Scenario: You need to provide Layer 7 routing capabilities for your microservices, including host-based routing and path-based routing. How do you implement this in Kubernetes?

Answer:

1. **Ingress Resource:** Use an **Ingress resource** to implement host-based and path-based routing.

- Example Ingress:

apiVersion: networking.k8s.io/v1

kind: Ingress

metadata:

name: microservice-ingress

spec:

rules:

- host: myapp.example.com

http:

paths:

- path: /service1

backend:

service:

name: service1

port:

number: 80

- path: /service2

backend:

service:

name: service2

port:

number: 80

Kubernetes Deployment Strategies

175. Scenario: You want to perform a rolling update of your application without downtime. How do you configure this in Kubernetes?

Answer:

1. **RollingUpdate Strategy:** Use the RollingUpdate strategy in your Deployment spec to ensure a smooth update without downtime.

- Example Deployment with rolling update strategy:

apiVersion: apps/v1

kind: Deployment

metadata:

name: myapp

spec:

replicas: 4

strategy:

type: RollingUpdate

rollingUpdate:

maxUnavailable: 1

maxSurge: 1

template:

spec:

containers:

- name: myapp-container

image: myapp:v2

176. Scenario: You want to implement a blue-green deployment for your application in Kubernetes. How do you configure this?

Answer:

1. **Create Two Deployments:** Create separate deployments for the blue and green versions of your application.

- Example blue deployment:

apiVersion: apps/v1

kind: Deployment

metadata:

name: blue-deployment

spec:

replicas: 3

template:

spec:

containers:

- name: myapp

image: myapp:v1

- Similarly, create the green deployment.

2. **Switch Service Selector:** When ready to switch traffic to the new version, update the Service selector to point to the green deployment.

Kubernetes - Advanced Troubleshooting

177. Scenario: Your pods are unable to pull images from a private Docker registry. How do you troubleshoot this issue?

Answer:

1. **Check Image Pull Secret:** Ensure that the **imagePullSecret** is configured correctly for the private registry.

- Example:

spec:

imagePullSecrets:

- name: regcred

- Create the secret:

bash

Copy code

```
kubectl create secret docker-registry regcred --docker-server=<registry-server> --  
docker-username=<username> --docker-password=<password>
```

178. Scenario: A pod is stuck in Terminating state and won't delete. How do you forcefully remove it?

Answer:

1. **Force Delete:** Use the following command to forcefully delete the pod:

```
kubectl delete pod <pod-name> --grace-period=0 --force
```

Kubernetes API Server and Control Plane

179. Scenario: The API server is experiencing high latency, and some requests are timing out. How do you troubleshoot and resolve this?

Answer:

1. **Check API Server Logs:** Review the API server logs for any errors or bottlenecks.
2. **Monitor Resource Utilization:** Use `kubectl top` to monitor the CPU and memory usage of the control plane components.
3. **Horizontal Scaling:** If necessary, scale the control plane components (e.g., API server) to handle more requests.

180. Scenario: You want to restrict the Kubernetes API server access to specific IP ranges. How do you implement this?

Answer:

1. **Firewall Rules:** Configure firewall rules or **security groups** (for cloud environments) to allow access to the API server only from trusted IP ranges.
 2. **API Server Flags:** Use the `--client-ca-file` flag to restrict access to clients with valid certificates and control who can connect to the API server.
-

Kubernetes - Storage Troubleshooting

181. Scenario: Your PersistentVolumeClaim (PVC) is stuck in Pending state. How do you troubleshoot this?

Answer:

1. **Check Events:** Run `kubectl describe pvc <pvc-name>` to check for events related to storage provisioning.
2. **StorageClass:** Ensure that the correct StorageClass is specified, and that the provisioner is functioning properly.
3. **PersistentVolume Availability:** Check if a matching PersistentVolume (PV) is available for the claim to bind.

182. Scenario: You need to resize a PersistentVolumeClaim (PVC) that has been dynamically provisioned. How do you do this?

Answer:

1. **Enable Volume Expansion:** Ensure that the StorageClass supports volume expansion by setting `allowVolumeExpansion: true`.
2. **Update PVC Size:** Modify the PVC to request additional storage:

`apiVersion: v1`

`kind: PersistentVolumeClaim`

`spec:`

`resources:`

`requests:`

`storage: 20Gi`

3. **Verify Resize:** Use `kubectl describe pvc <pvc-name>` to verify that the PVC has been resized.

Kubernetes Advanced Security

183. Scenario: You need to implement pod-level security policies to prevent privileged containers from running in your cluster. How do you configure this?

Answer:

1. **PodSecurityPolicies (PSP):** Create a PodSecurityPolicy that prevents privileged containers from running.

- Example PSP:

`apiVersion: policy/v1beta1`

`kind: PodSecurityPolicy`

`metadata:`

`name: restricted-psp`

spec:

privileged: false

runAsUser:

rule: MustRunAsNonRoot

184. Scenario: You need to audit API access to your Kubernetes cluster for security compliance. How do you enable Kubernetes audit logging?

Answer:

1. **Enable Audit Logging:** Configure the API server to enable audit logging by setting the --audit-policy-file and --audit-log-path flags.

- Example audit policy:

apiVersion: audit.k8s.io/v1

kind: Policy

rules:

- level: RequestResponse

resources:

- group: ""

resources: ["pods", "services"]

2. **Centralize Logs:** Use a centralized logging system like Elasticsearch or Splunk to store and analyze audit logs.

Kubernetes - Advanced Scheduling

185. Scenario: You need to ensure that a specific workload is scheduled on nodes with high CPU capacity. How do you configure this in Kubernetes?

Answer:

1. **Node Affinity:** Use **Node Affinity** to ensure that the workload is scheduled on nodes with the appropriate capacity.

- Example node affinity configuration:

spec:

affinity:

nodeAffinity:

requiredDuringSchedulingIgnoredDuringExecution:

nodeSelectorTerms:

- matchExpressions:

- key: kubernetes.io/instance-type

operator: In

values:

- high-cpu

186. Scenario: You want to ensure that a workload runs only on nodes in a specific availability zone. How do you enforce this scheduling constraint?

Answer:

1. **Node Affinity:** Use **Node Affinity** to restrict scheduling to nodes in a specific availability zone.

- Example:

spec:

affinity:

nodeAffinity:

requiredDuringSchedulingIgnoredDuringExecution:

nodeSelectorTerms:

- matchExpressions:

- key: topology.kubernetes.io/zone

operator: In

values:

- us-west-1a

Kubernetes Networking

187. Scenario: You need to restrict external access to a Kubernetes service, allowing traffic only from internal cluster components. How do you configure this?

Answer:

1. **ClusterIP Service:** Use a **ClusterIP** service to restrict access to within the cluster.
2. **Network Policy:** Use a **Network Policy** to explicitly deny all external traffic:

apiVersion: networking.k8s.io/v1

kind: NetworkPolicy

metadata:

name: internal-only-policy

spec:

podSelector: {}

policyTypes:

- Ingress

ingress:

- from:

- podSelector: {}

Kubernetes Debugging

188. Scenario: You need to troubleshoot a container that has crashed but need to examine the state of the filesystem after the crash. How do you do this?

Answer:

1. **Ephemeral Containers:** Use **ephemeral containers** to attach to a running pod and examine the container's filesystem:

```
kubectrl debug -it <pod-name> --image=busybox --target=<container-name>
```

189. Scenario: You notice network performance degradation between pods. How do you investigate the issue?

Answer:

1. **Network Performance Tools:** Use tools like **Weave Scope** or **Prometheus** to monitor network traffic between pods.
2. **kube-proxy Logs:** Check the **kube-proxy** logs for any issues with packet forwarding or iptables rules.

Kubernetes Custom Controllers

190. Scenario: You need to implement custom Kubernetes controllers to automate specific actions (e.g., scaling based on external events). How do you create a custom controller?

Answer:

1. **Kubebuilder:** Use **Kubebuilder** to create a custom controller.

- Initialize the project:

```
kubebuilder init --domain mycompany.com
```

- Create an API and controller:

```
kubebuilder create api --group mygroup --version v1 --kind MyController
```

2. **Controller Logic:** Implement custom logic in the controller for specific actions such as scaling or updating resources.

191. Scenario: You need to implement a custom resource definition (CRD) for managing a custom resource type in your Kubernetes cluster. How do you create this CRD?

Answer:

1. **Define CRD:** Create a CRD YAML file that defines your custom resource.

- Example CRD:

apiVersion: apiextensions.k8s.io/v1

kind: CustomResourceDefinition

metadata:

name: myresources.mygroup.mycompany.com

spec:

group: mygroup.mycompany.com

versions:

- name: v1

served: true

storage: true

scope: Namespaced

names:

plural: myresources

singular: myresource

kind: MyResource

Kubernetes Advanced Networking

192. Scenario: You need to configure multi-tenant networking, ensuring that each tenant's services are isolated. How do you achieve this in Kubernetes?

Answer:

1. **Network Policies:** Use **Network Policies** to isolate tenant services within their own namespaces.

- Example:

apiVersion: networking.k8s.io/v1

kind: NetworkPolicy

metadata:

name: tenant-a-policy

namespace: tenant-a

spec:

podSelector:

matchLabels:

app: tenant-a-app

policyTypes:

- Ingress

ingress:

- from:

- podSelector: {}

Kubernetes Secret Management

193. Scenario: You need to securely manage and rotate secrets used by your Kubernetes applications. How do you implement this?

Answer:

1. **HashiCorp Vault:** Use **HashiCorp Vault** to manage and inject secrets dynamically into pods.

- Example of injecting secrets via Vault:

```
vault kv put secret/myapp/api-key value=mysecretvalue
```

2. **Kubernetes Secrets:** Alternatively, use **Kubernetes Secrets** and mount them as environment variables in pods:

env:

```
- name: API_KEY
```

```
valueFrom:
```

```
secretKeyRef:
```

```
  name: api-secret
```

```
  key: api-key
```

Kubernetes Custom Resources

194. Scenario: You need to create a custom resource to manage a new type of resource in your Kubernetes environment. How do you define and implement this custom resource?

Answer:

1. **Custom Resource Definition (CRD):** Create a CRD that defines the new resource type.

- Example:

```
apiVersion: apiextensions.k8s.io/v1
```

```
kind: CustomResourceDefinition
```

```
metadata:
```

```
  name: customresources.mycompany.com
```

spec:

group: mycompany.com

versions:

- name: v1

served: true

storage: true

names:

plural: customresources

singular: customresource

kind: CustomResource

2. **Custom Controller:** Implement a custom controller using **Kubebuilder** to manage the lifecycle of the custom resource.

Kubernetes Security Best Practices

195. Scenario: Your organization wants to enforce image scanning for vulnerabilities before allowing images to run in production. How do you implement this in Kubernetes?

Answer:

1. **Trivy or Clair:** Integrate image vulnerability scanners such as **Trivy** or **Clair** into your CI/CD pipeline.
 - o Example:

trivy image --exit-code 1 myapp:latest

2. **Admission Controller:** Use an admission controller to enforce image policies. **OPA Gatekeeper** can be configured to block images with vulnerabilities.

Kubernetes Scheduling - Advanced

196. Scenario: You need to schedule a workload on nodes that have GPUs available. How do you ensure this in Kubernetes?

Answer:

1. **Use GPU Resources:** Specify the `nvidia.com/gpu` resource in the pod's resource requests.

- Example:

resources:

limits:

`nvidia.com/gpu: 1`

197. Scenario: You need to ensure a critical application always runs on the most powerful nodes in your cluster. How do you enforce this scheduling policy?

Answer:

1. **Node Affinity:** Use node affinity to schedule the application on nodes with specific hardware capabilities (e.g., high memory, high CPU).

spec:

affinity:

nodeAffinity:

`requiredDuringSchedulingIgnoredDuringExecution:`

`nodeSelectorTerms:`

- `matchExpressions:`

- `key: kubernetes.io/instance-type`

- `operator: In`

- `values:`

- `high-memory`

Kubernetes Cost Optimization

198. Scenario: Your Kubernetes cluster is over-provisioned, leading to higher infrastructure costs. How do you optimize the resources and reduce costs?

Answer:

1. **Use Resource Requests and Limits:** Ensure that all pods have resource requests and limits defined, avoiding over-provisioning.
 2. **Vertical Pod Autoscaler (VPA):** Use VPA to automatically adjust resource requests based on actual usage, preventing over-allocation.
 3. **Cluster Autoscaler:** Enable **Cluster Autoscaler** to remove underutilized nodes from the cluster.
-

Kubernetes - Advanced HA and Fault Tolerance

199. Scenario: You need to ensure that your Kubernetes cluster control plane is highly available. How do you configure this?

Answer:

1. **Multiple Masters:** Set up multiple master nodes across different availability zones, using a load balancer in front of the API server.
2. **HA etcd:** Set up a highly available **etcd** cluster with an odd number of members (e.g., 3 or 5) to ensure quorum-based consensus.

200. Scenario: You need to ensure that workloads remain highly available even during node failures. How do you implement this in Kubernetes?

Answer:

1. **Pod Disruption Budgets (PDB):** Use PDBs to ensure that a minimum number of replicas are always available:

apiVersion: policy/v1

kind: PodDisruptionBudget

metadata:

name: myapp-pdb

spec:

minAvailable: 80%

selector:

matchLabels:

app: myapp

2. **Cluster Autoscaler:** Use **Cluster Autoscaler** to automatically replace failed nodes.